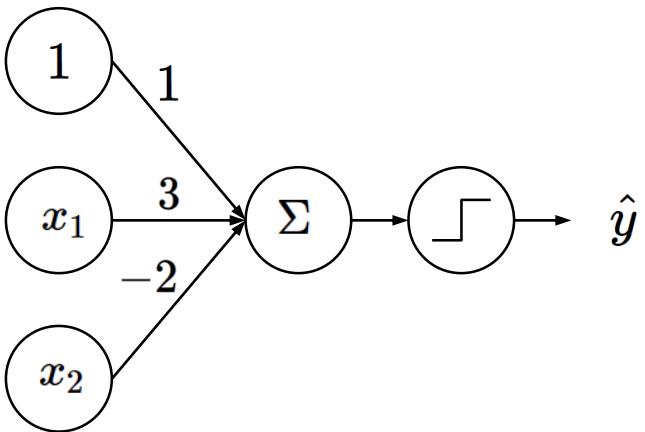




Artificial Neural Networks

**Prof. Seungchul Lee
Industrial AI Lab.**

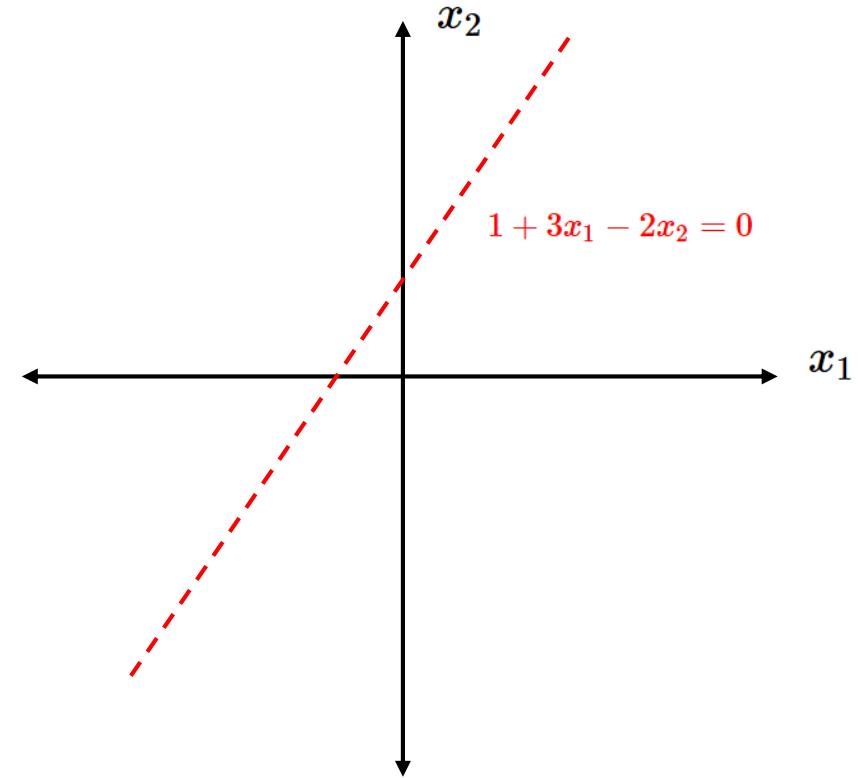
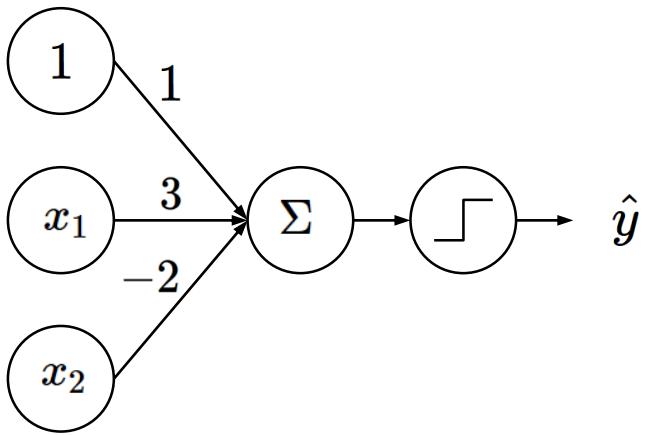
Perceptron: Example



$$\begin{aligned}\hat{y} &= g(\omega_0 + X^T \omega) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2)\end{aligned}$$

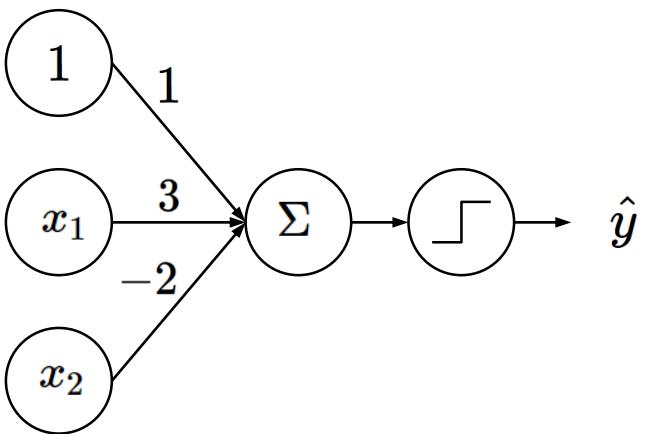
Perceptron: Example

$$\hat{y} = g(1 + 3x_1 - 2x_2)$$

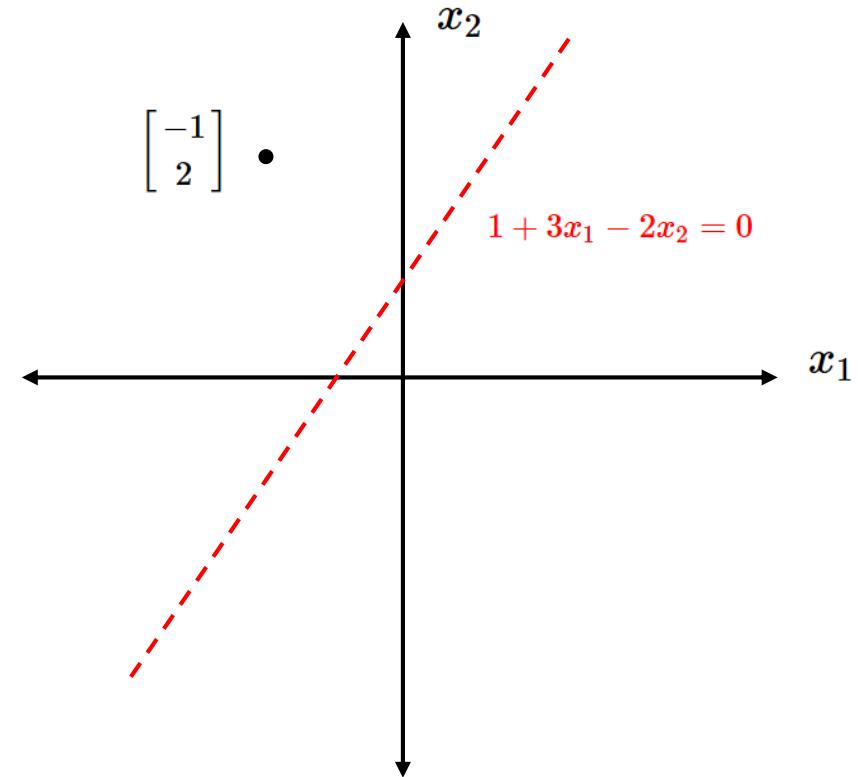


Perceptron: Example

$$\hat{y} = g(1 + 3x_1 - 2x_2)$$

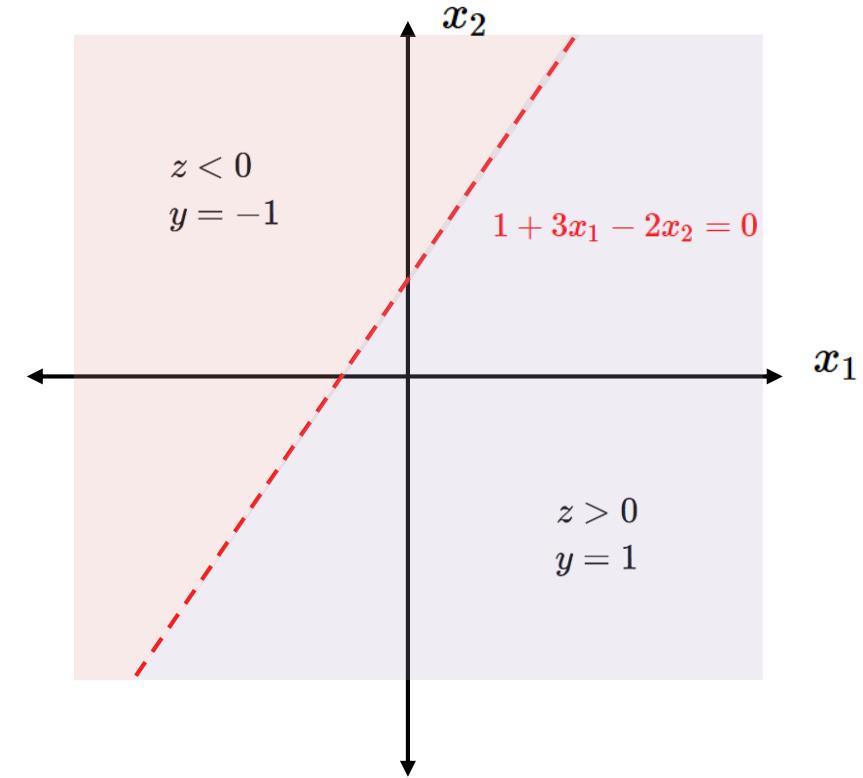
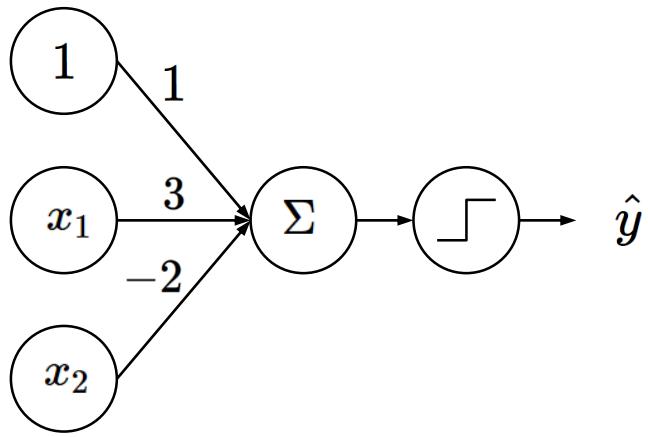


$$\hat{y} = g(1 + 3 \times (-1) - 2 \times 2) = g(-6) = -1$$

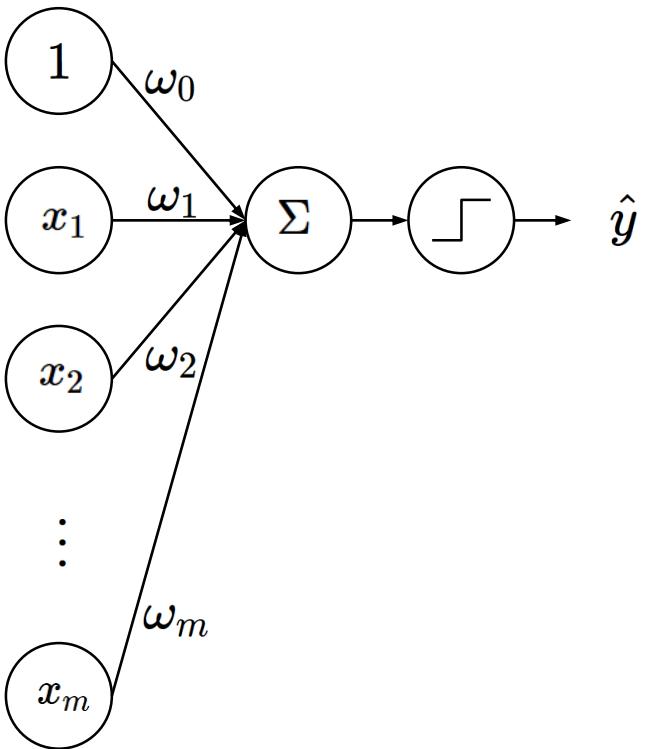


Perceptron: Example

$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Perceptron: Forward Propagation



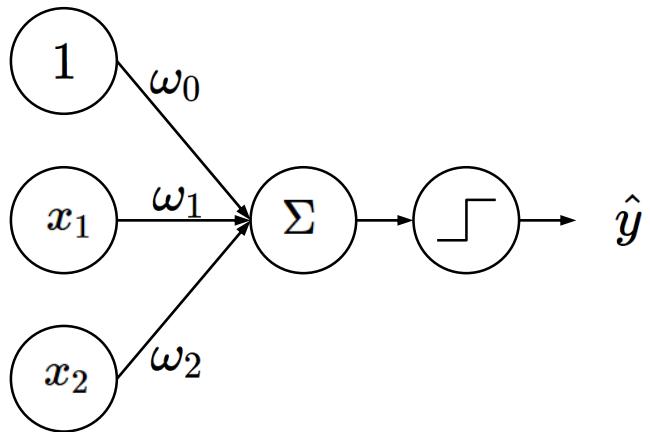
$$\hat{y} = g(\omega_0 + X^T \omega)$$

$$= g\left(\omega_0 + \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}^T \begin{bmatrix} \omega_1 \\ \vdots \\ \omega_m \end{bmatrix}\right)$$

From Perceptron to MLP

Artificial Neural Networks: Perceptron

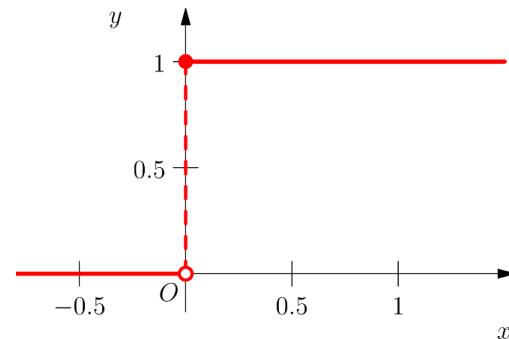
- Perceptron for $h(\theta)$ or $h(\omega)$
 - Neurons compute the weighted sum of their inputs
 - A neuron is activated or fired when the sum a is positive



- A step function is not differentiable
- One neuron is often not enough
 - One hyperplane

$$a = \omega_0 + \omega_1 x_1 + \omega_2 x_2$$

$$\hat{y} = g(a) = \begin{cases} 1 & a > 0 \\ 0 & \text{otherwise} \end{cases}$$

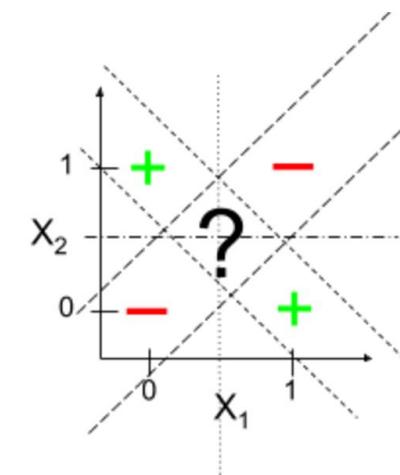
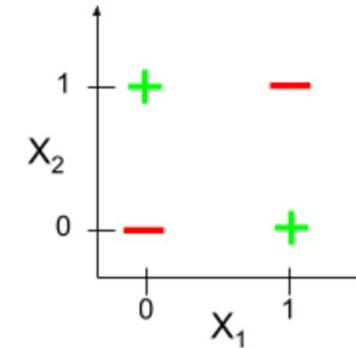
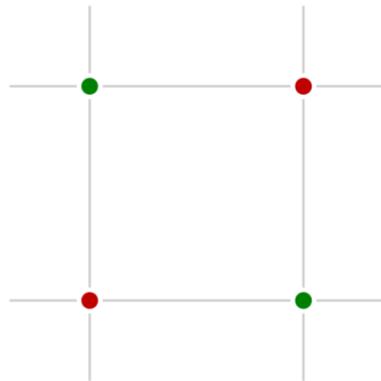


Here, a step function is illustrated instead of a sign function

XOR Problem

- Minsky-Papert Controversy on XOR
 - Not linearly separable
 - Limitation of perceptron

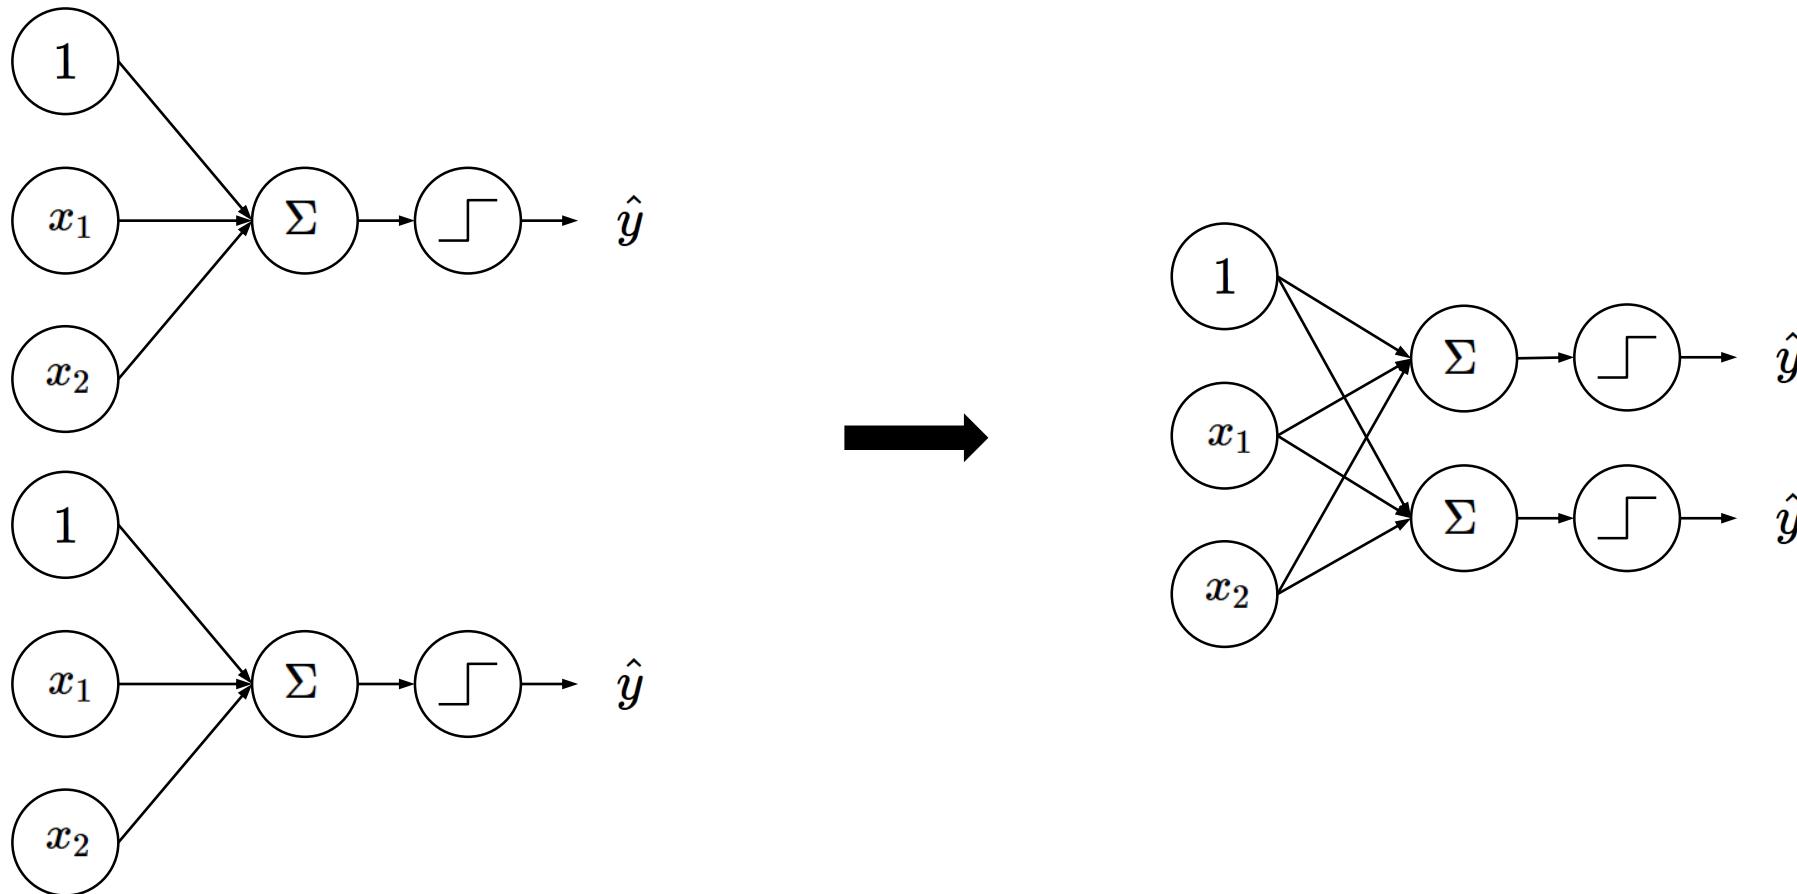
x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0



- Single neuron = one linear classification boundary

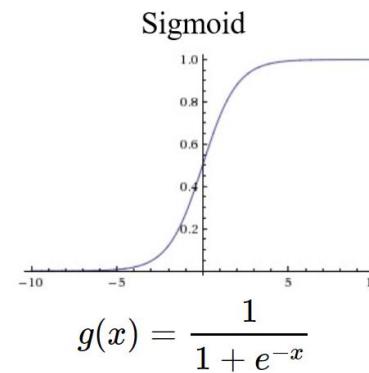
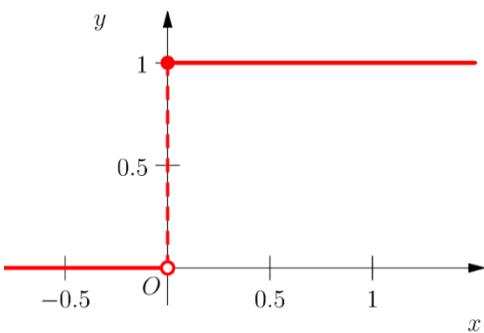
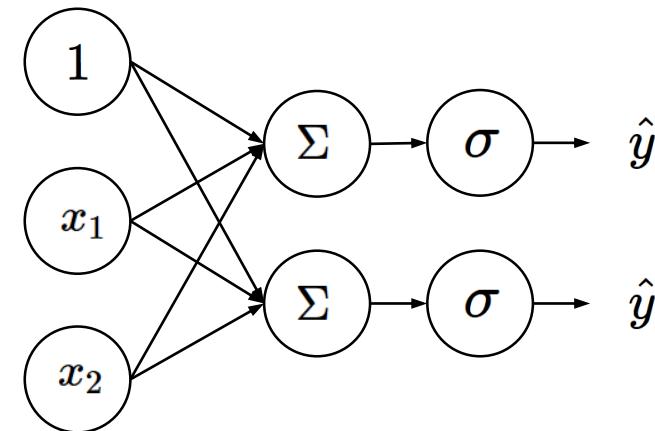
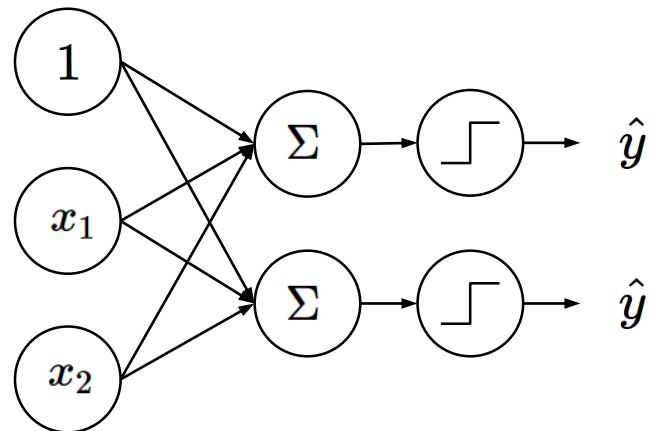
Artificial Neural Networks: MLP

- Multi-layer Perceptron (MLP) = Artificial Neural Networks (ANN)
 - Multi neurons = multiple linear classification boundaries



Artificial Neural Networks: Activation Function

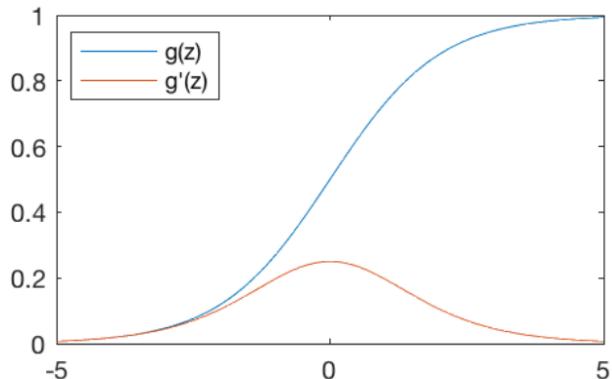
- Differentiable nonlinear activation function



Common Activation Functions

Discuss later

Sigmoid Function

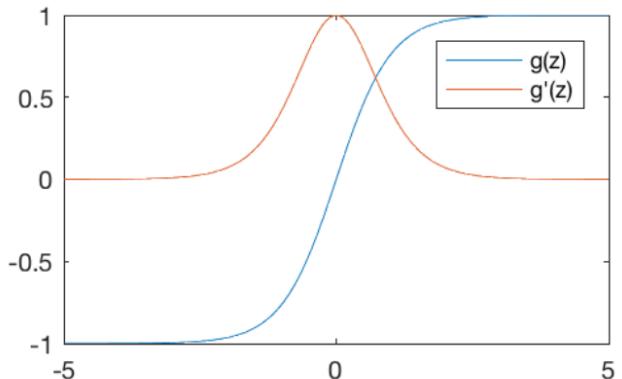


$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

 `tf.nn.sigmoid(z)`

Hyperbolic Tangent

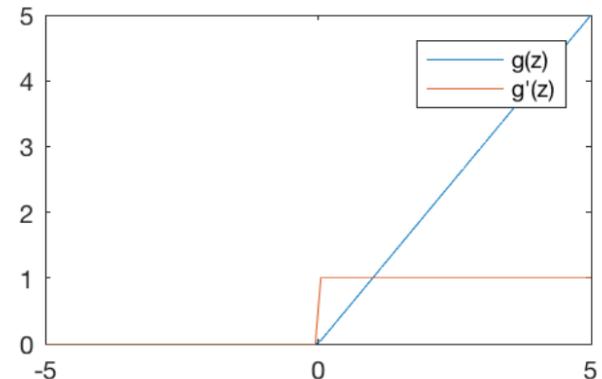


$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

 `tf.nn.tanh(z)`

Rectified Linear Unit (ReLU)



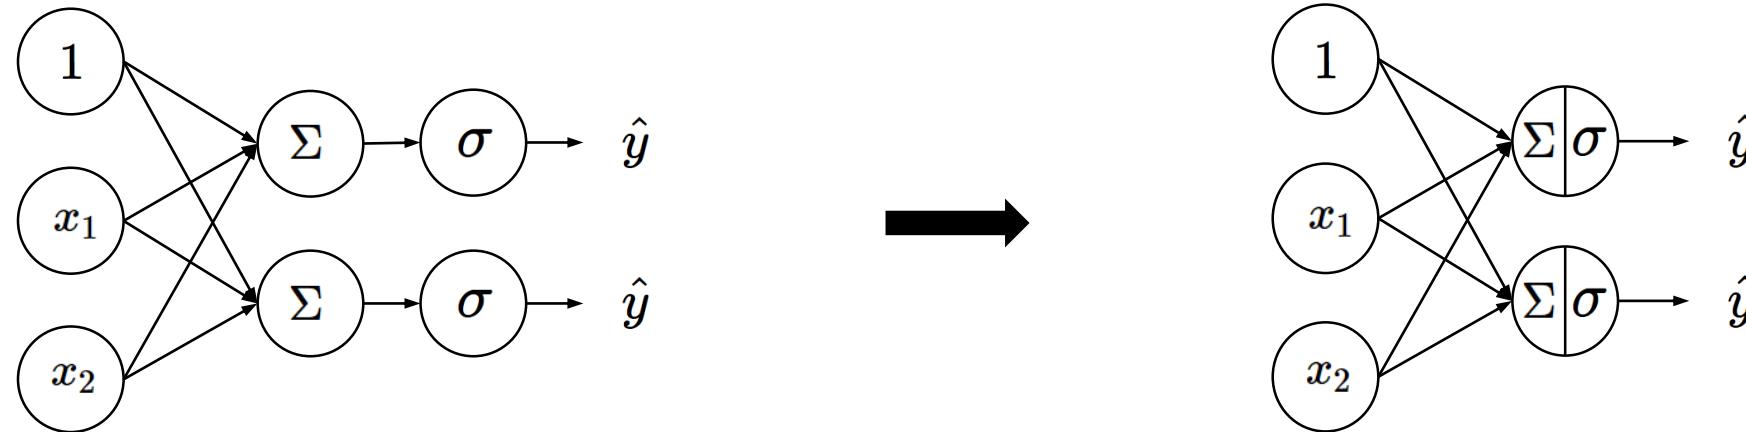
$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

 `tf.nn.relu(z)`

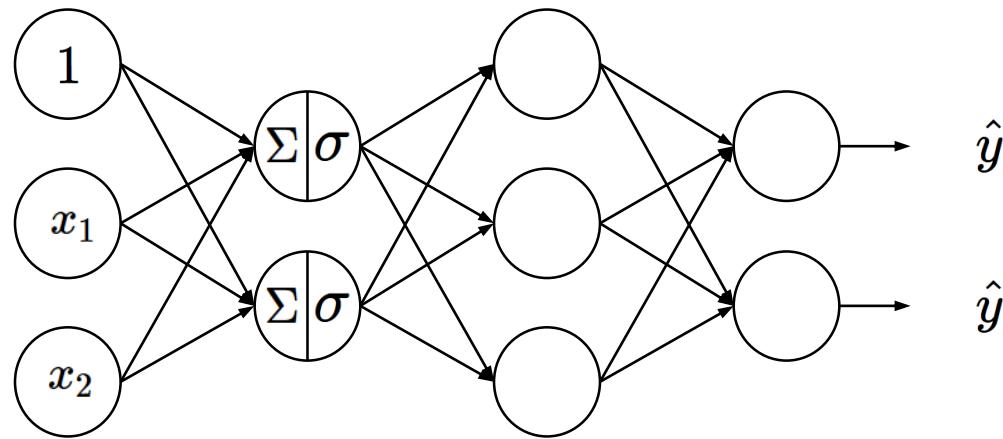
Artificial Neural Networks

- In a compact representation



Artificial Neural Networks

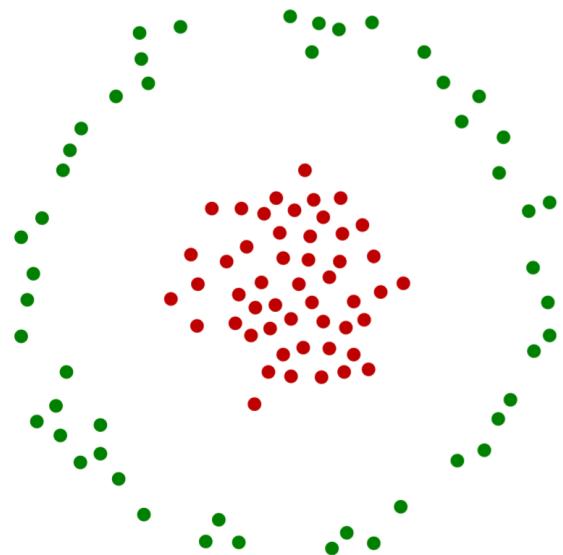
- A single layer is not enough to be able to represent complex relationship between input and output
⇒ perceptron with many layers and units



- Multi-layer perceptron
 - Features of features
 - Mapping of mappings

Another Perspective: ANN as Kernel Learning

Nonlinear Classification



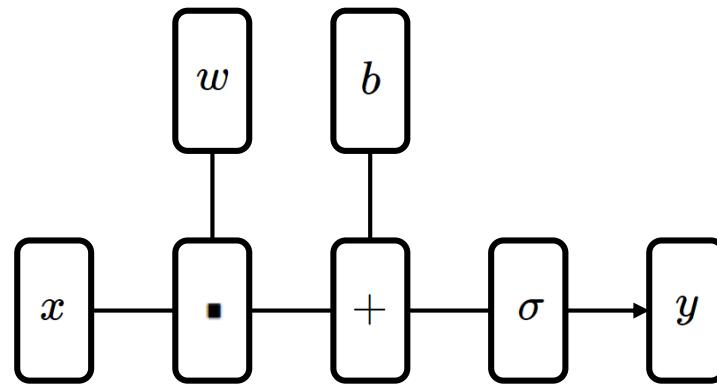
SVM with a polynomial
Kernel visualization

Created by:
Udi Aharoni

Neuron

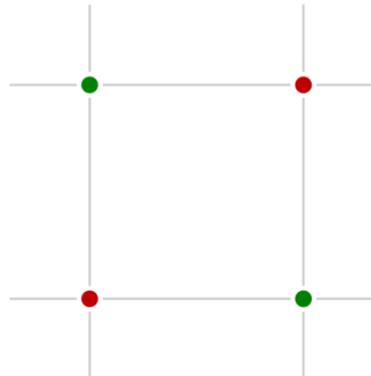
- We can represent this “neuron” as follows:

$$f(x) = \sigma(w \cdot x + b)$$



XOR Problem

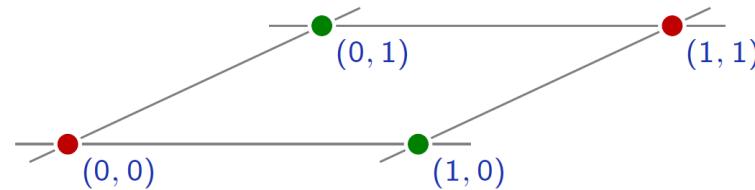
- The main weakness of linear predictors is their lack of capacity.
- For classification, the populations have to be linearly separable.



“xor”

Nonlinear Mapping

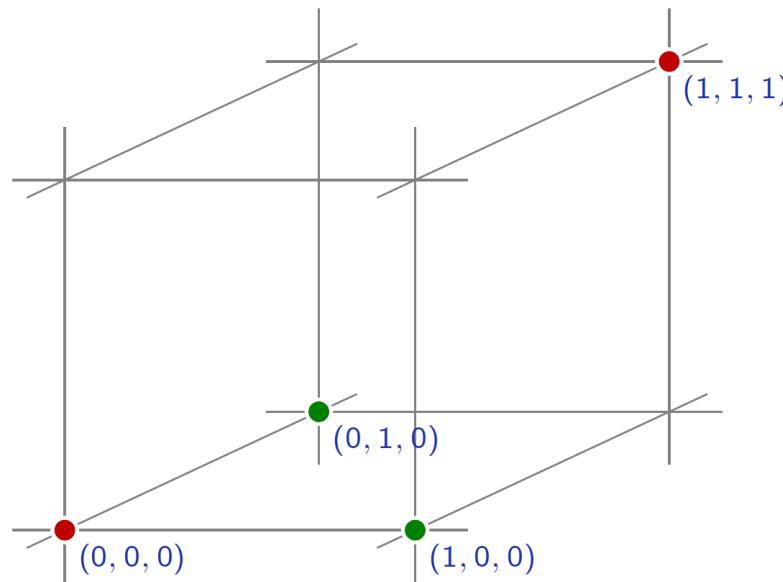
- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.

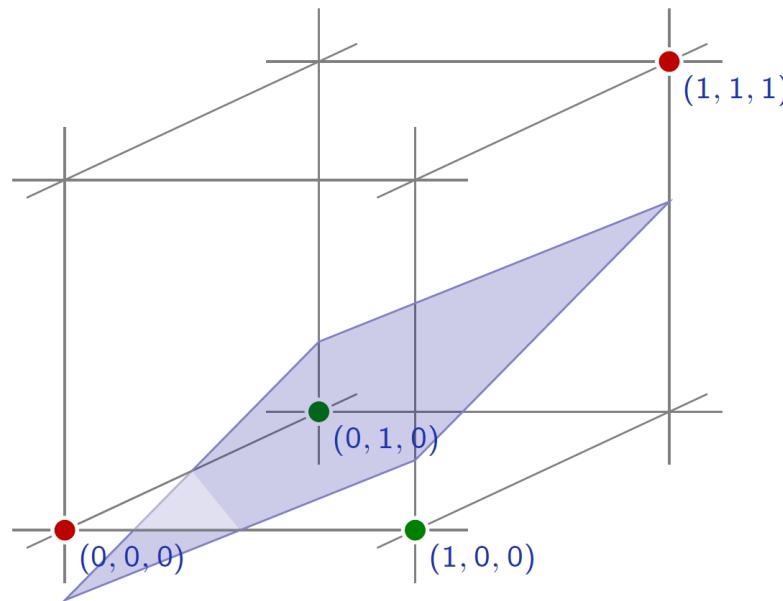
$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



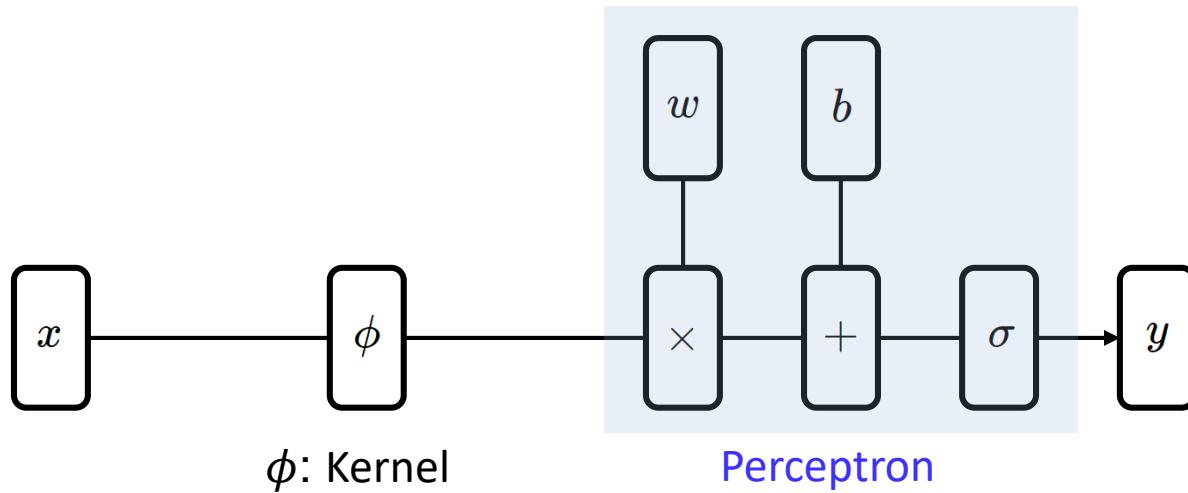
Kernel

- Often we want to capture nonlinear patterns in the data
 - nonlinear regression: input and output relationship may not be linear
 - nonlinear classification: classes may not be separable by a linear boundary
- Linear models (e.g. linear regression, linear SVM) are not just rich enough
 - by mapping data to higher dimensions where it exhibits linear patterns
 - apply the linear model in the new input feature space
 - mapping = changing the feature representation
- Kernels: make linear model work in nonlinear settings

Kernel + Neuron

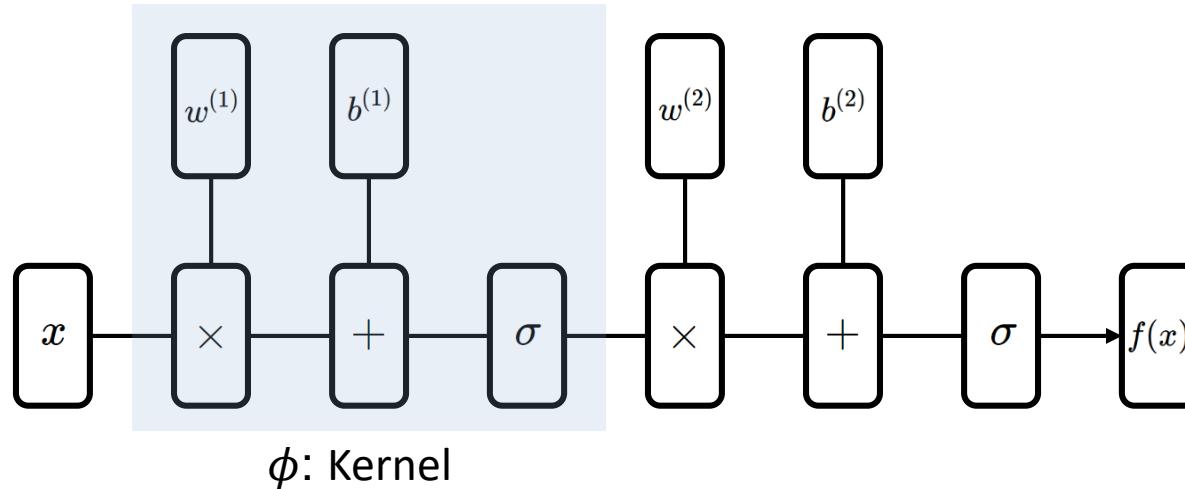
- Nonlinear mapping + neuron

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Neuron + Neuron

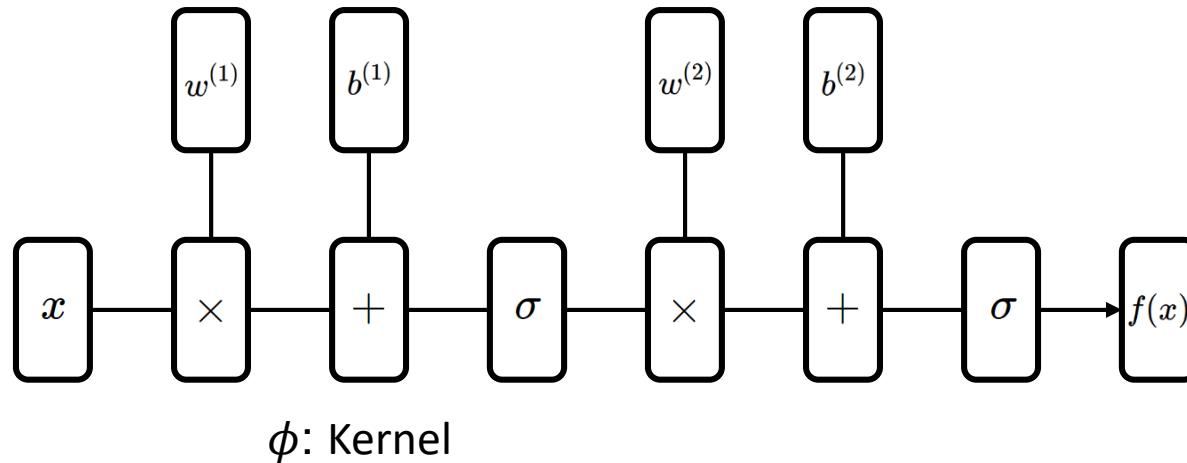
- Nonlinear mapping can be represented by another neurons



- Nonlinear Kernel
 - Nonlinear activation functions

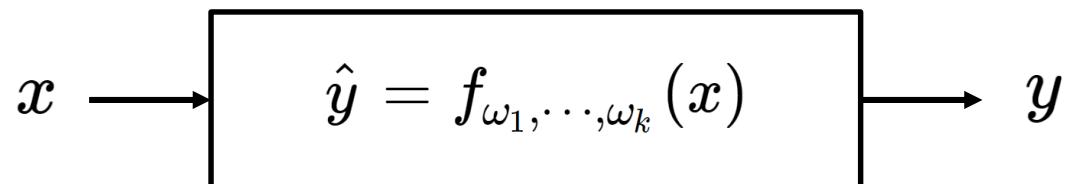
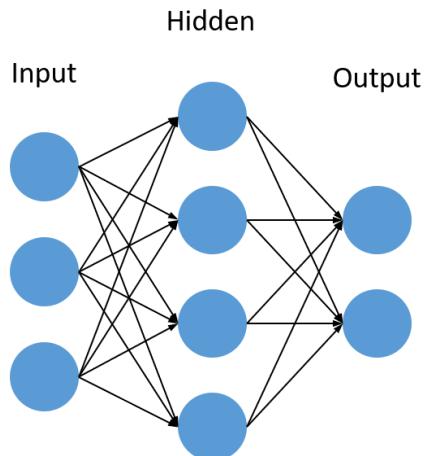
Multi Layer Perceptron

- Nonlinear mapping can be represented by another neurons
- We can generalize an MLP



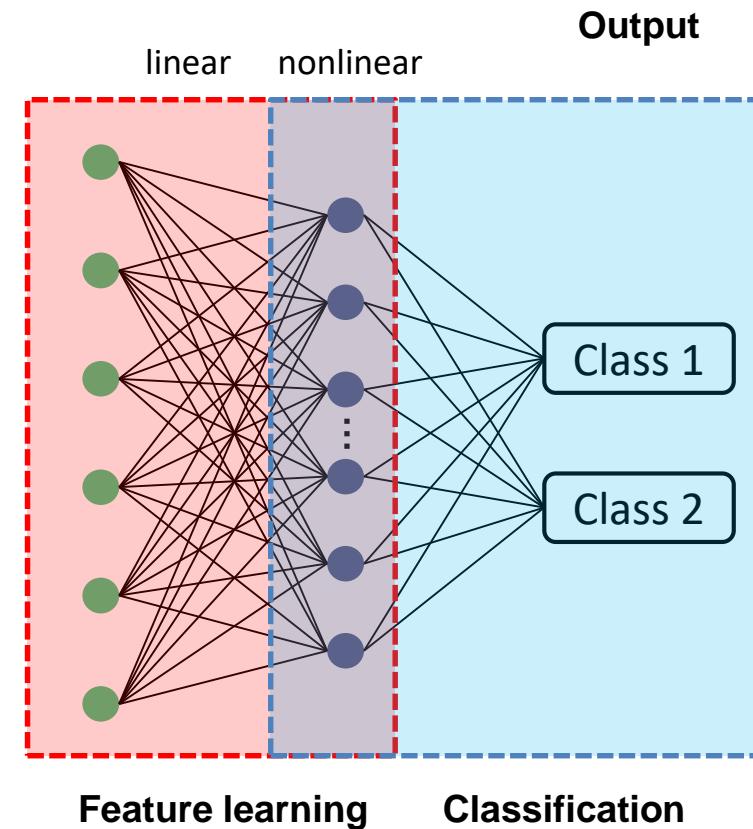
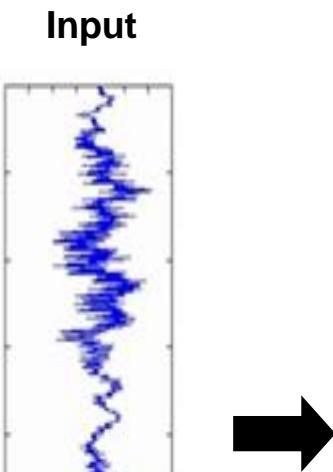
Summary

- Universal function approximator
- Universal function classifier
- Parameterized



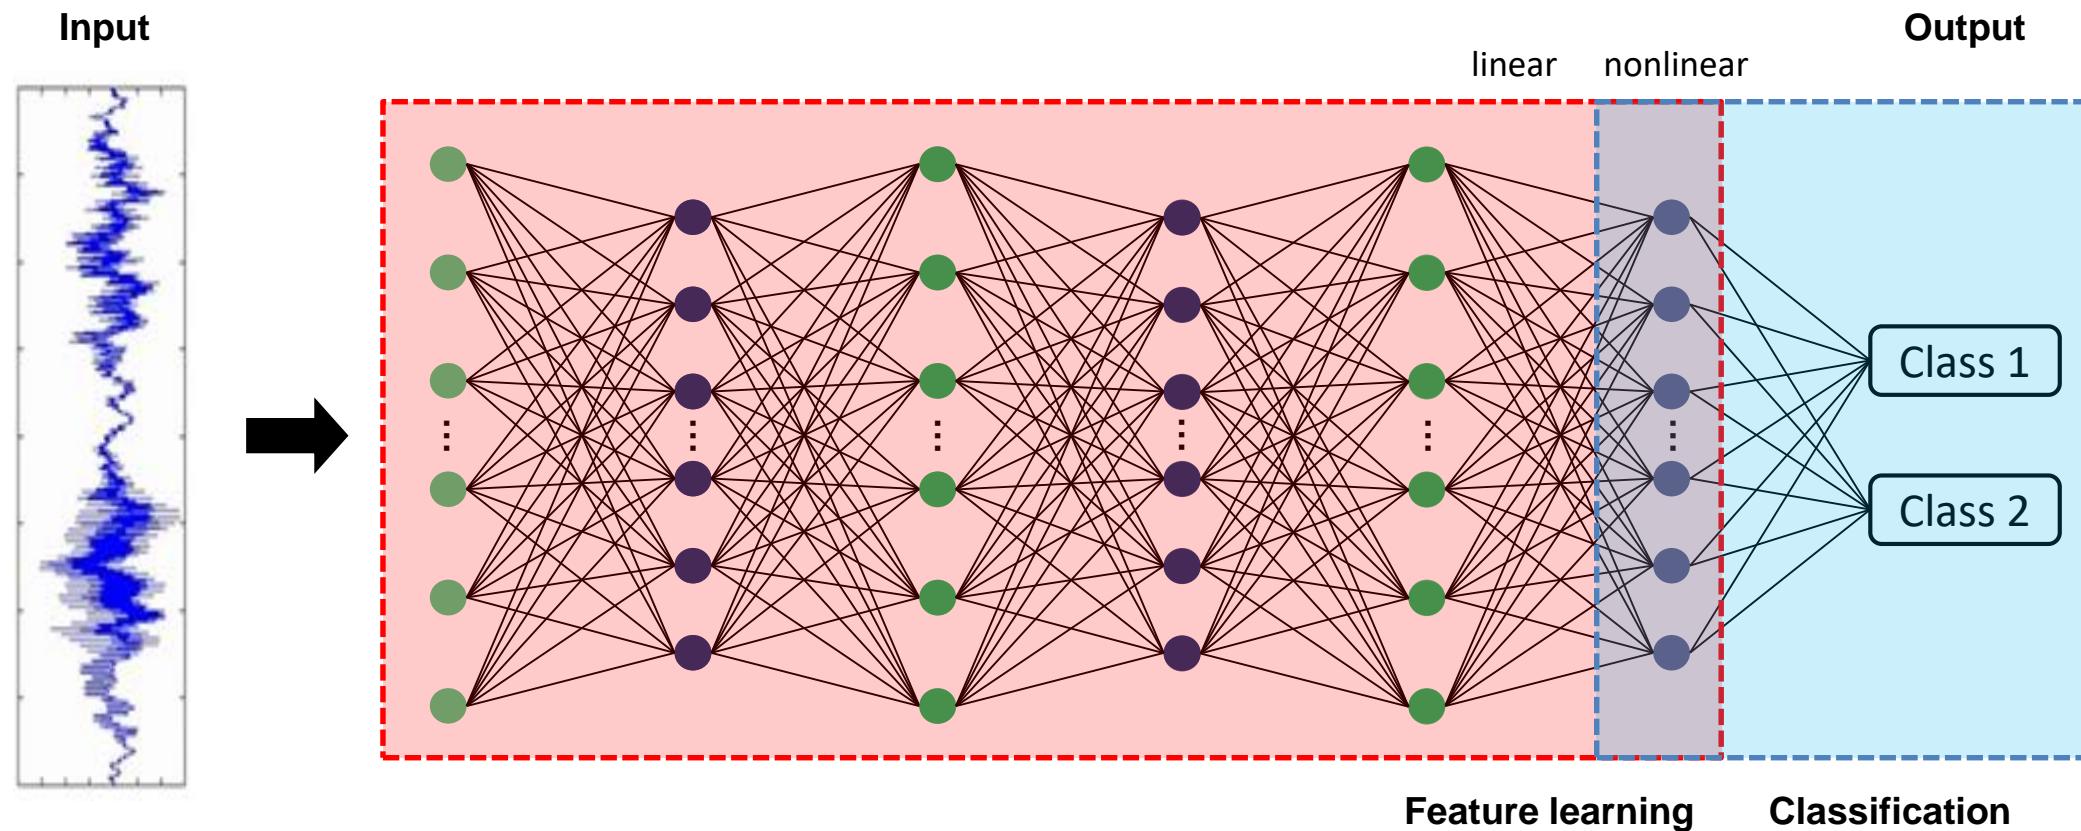
Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons



Deep Artificial Neural Networks

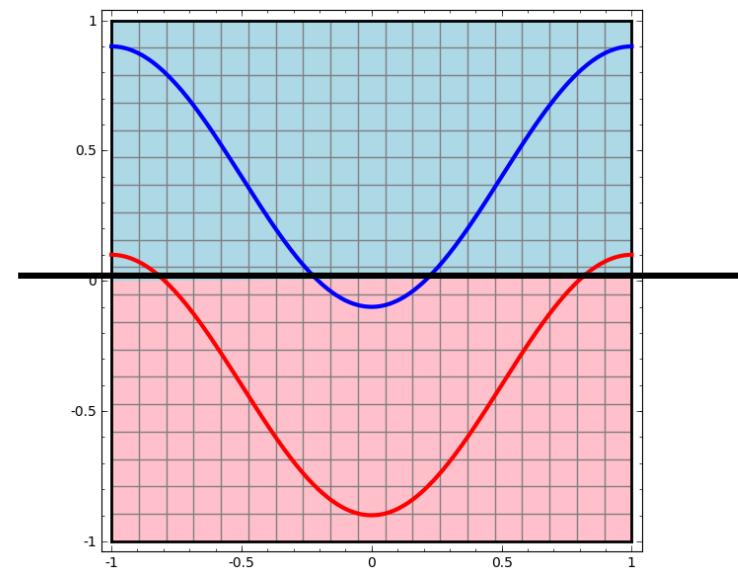
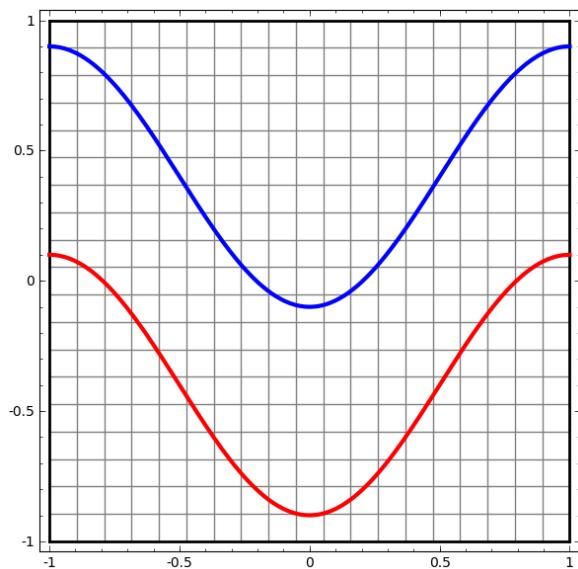
- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons



Looking at Parameters

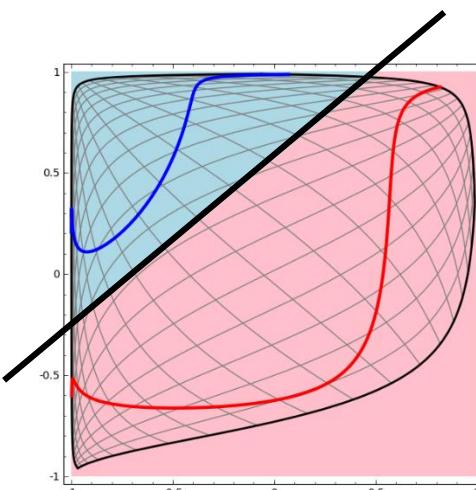
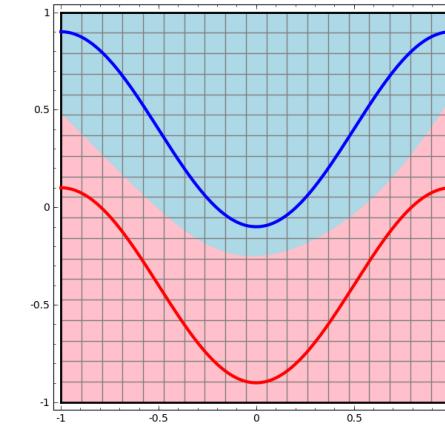
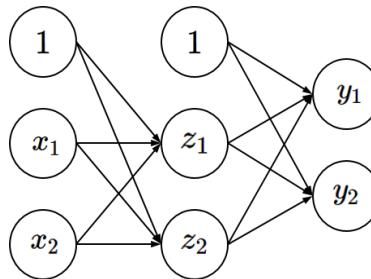
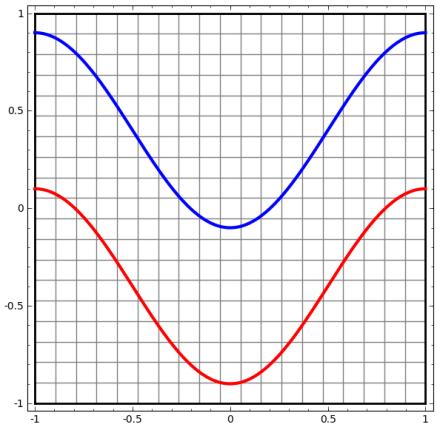
Example: Linear Classifier

- Perceptron tries to separate the two classes of data by dividing them with a line

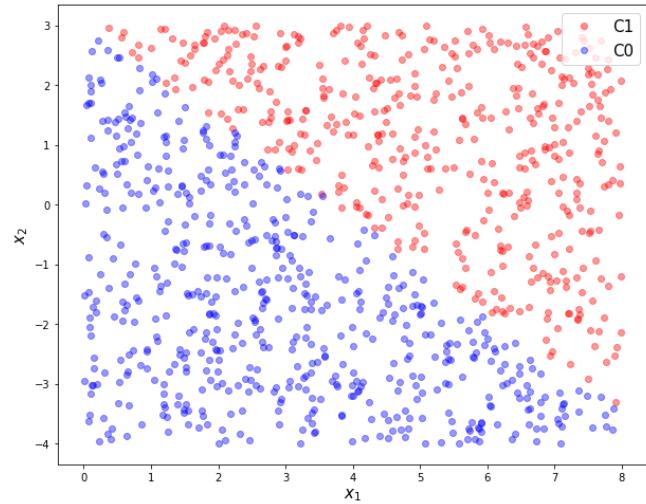


Example: Neural Networks

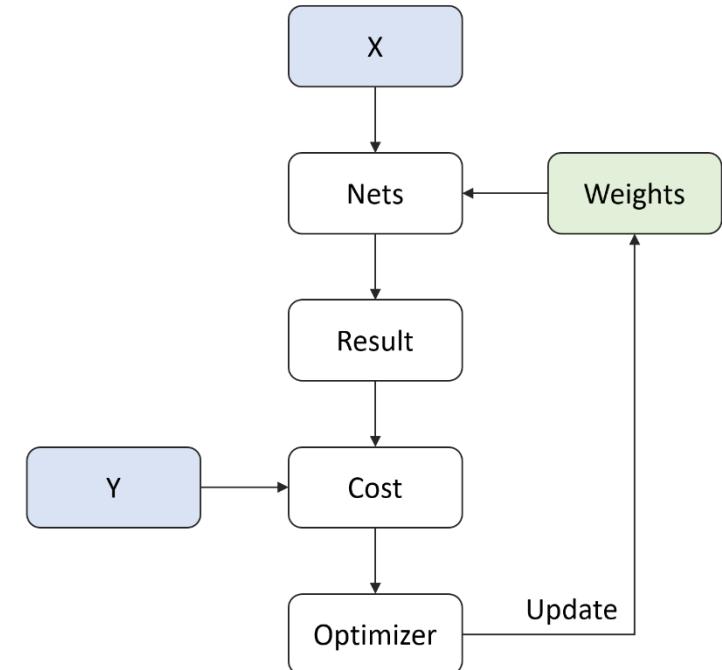
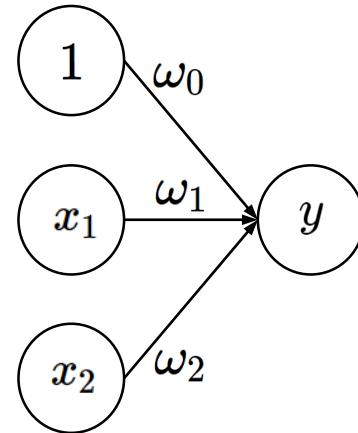
- The hidden layer learns a representation so that the data gets linearly separable



Logistic Regression in a Form of Neural Network



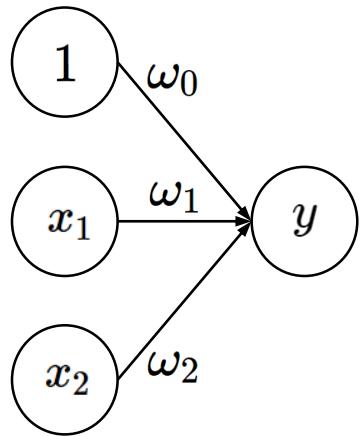
$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$



Logistic Regression in a Form of Neural Network

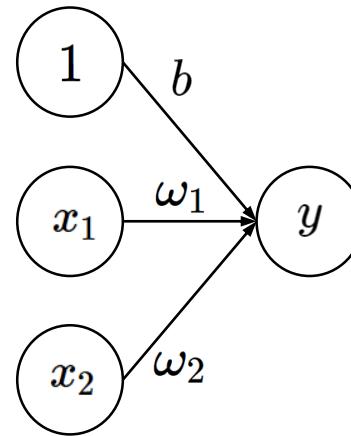
- Neural network convention

$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$

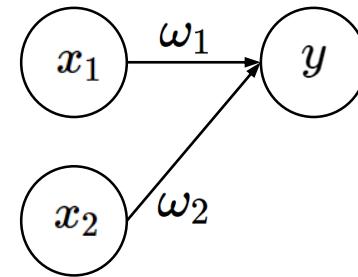


$\omega_0 \rightarrow b$

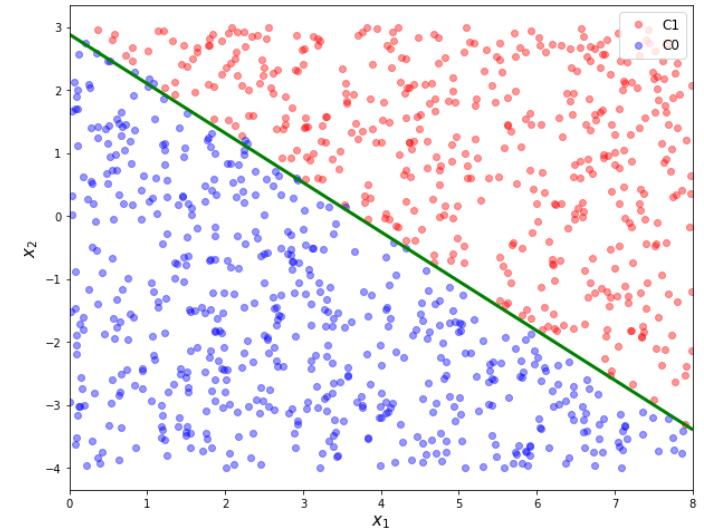
$$y = \sigma(b + \omega_1 x_1 + \omega_2 x_2)$$



Do not indicate bias units



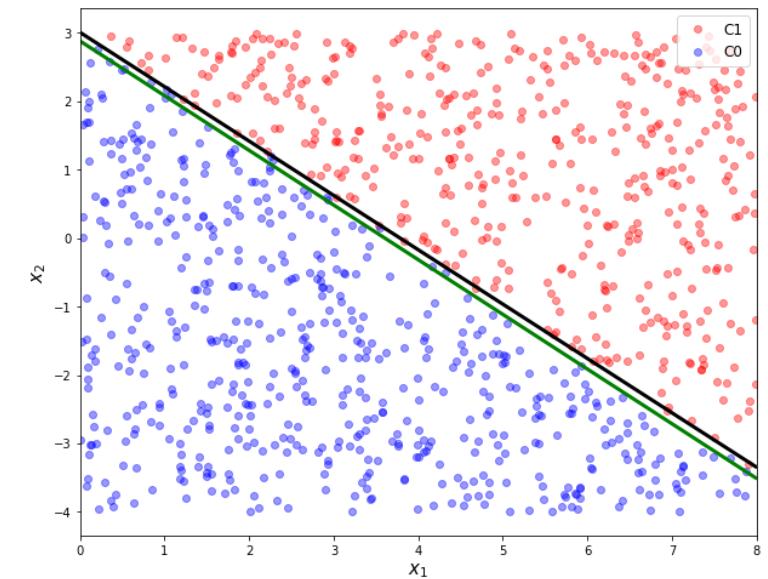
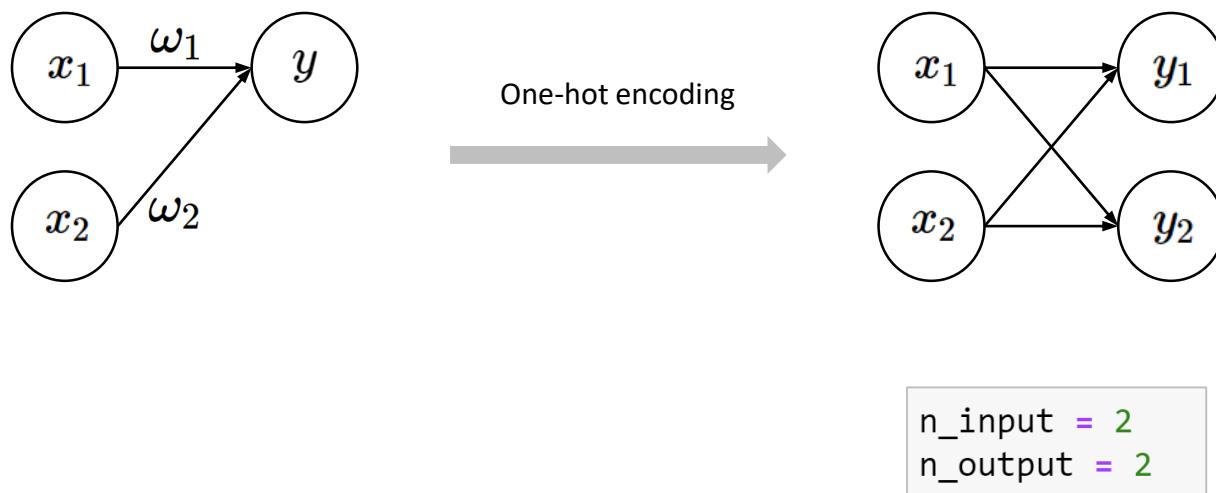
```
n_input = 2  
n_output = 1
```



Logistic Regression in a Form of Neural Network

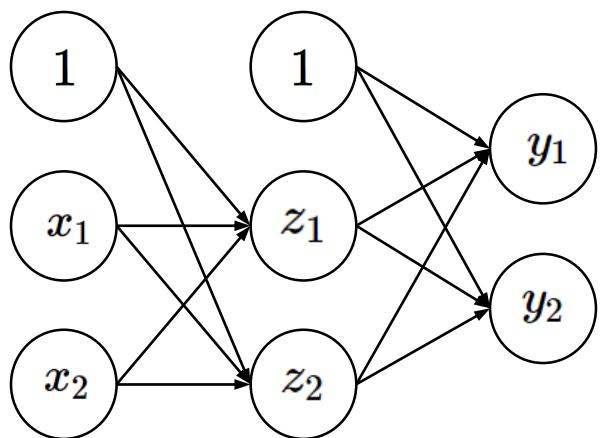
- One-hot encoding
 - One-hot encoding is a conventional practice for a multi-class classification

$$y^{(i)} \in \{1, 0\} \quad \Rightarrow \quad y^{(i)} \in \{[0, 1], [1, 0]\}$$

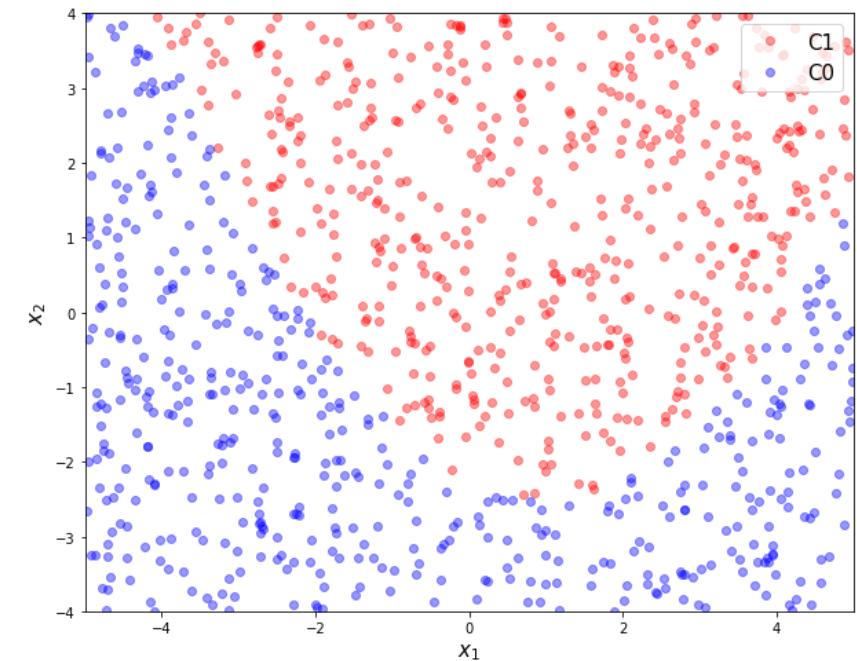
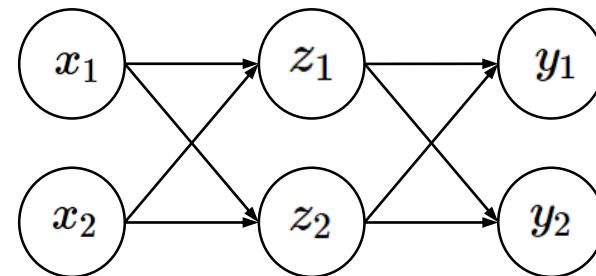


Nonlinearly Distributed Data

- Example to understand network's behavior
 - Include a hidden layer



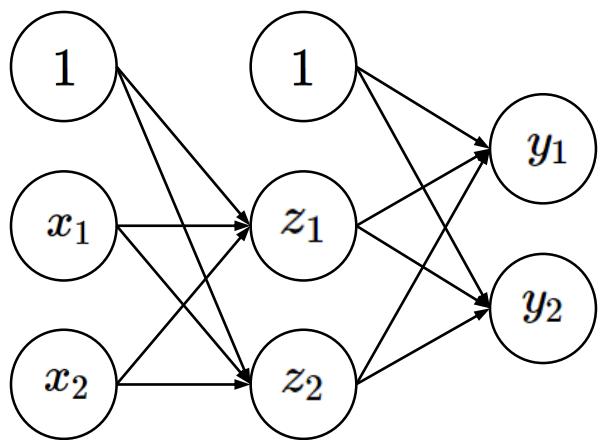
Do not include bias units



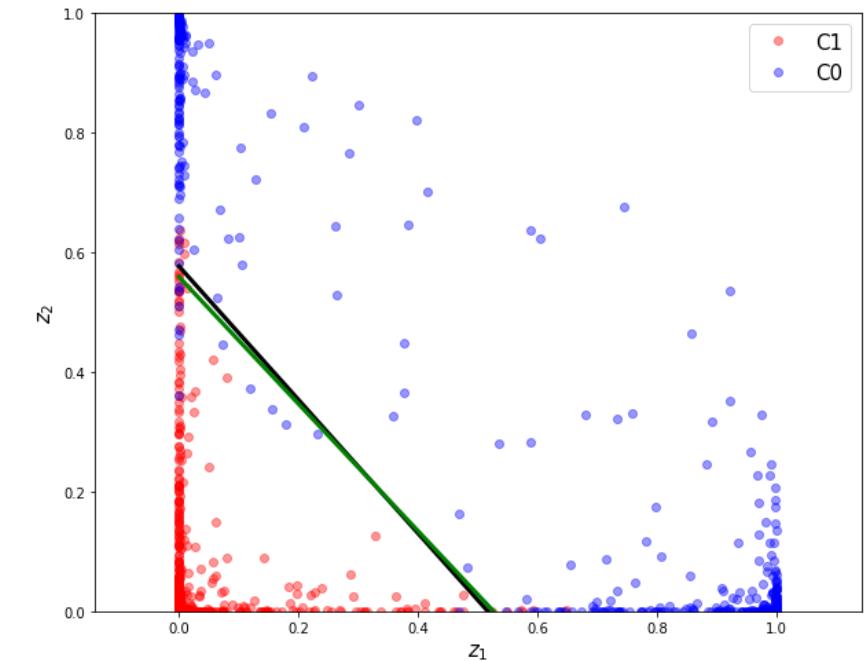
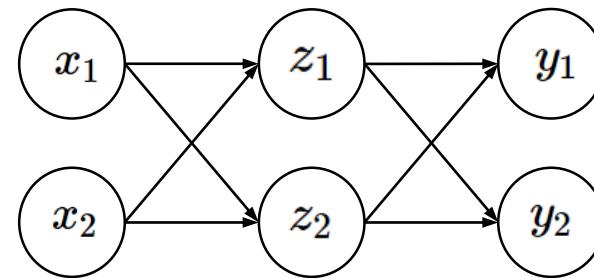
n_input = 2
n_hidden = 2
n_output = 2

Multi Layers

- z space



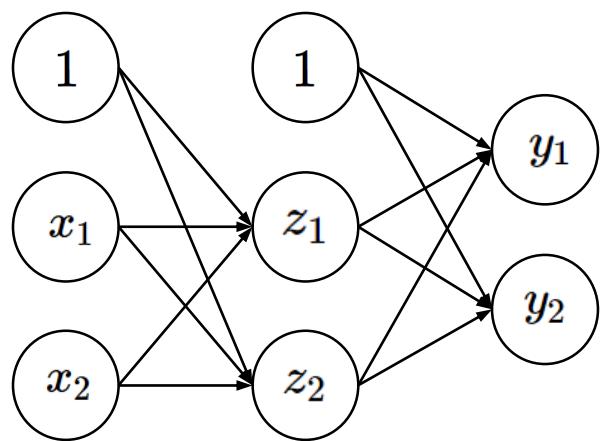
Do not include bias units



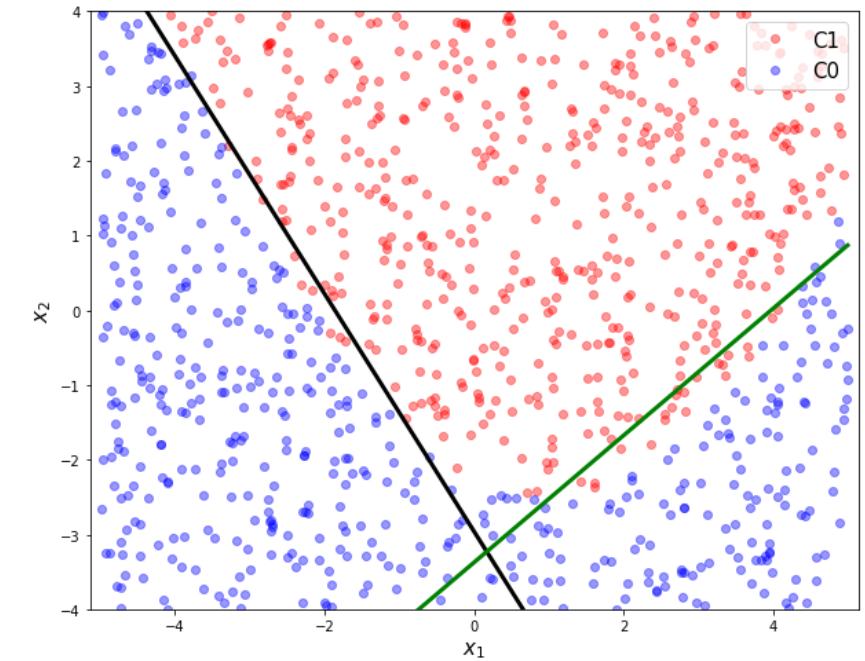
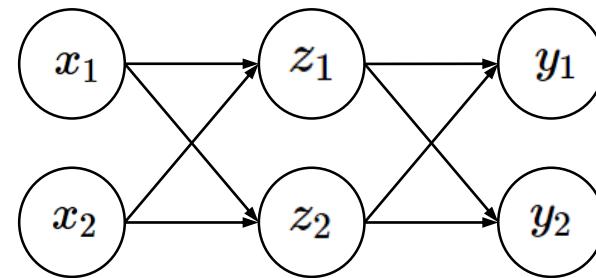
n_input = 2
n_hidden = 2
n_output = 2

Multi Layers

- x space



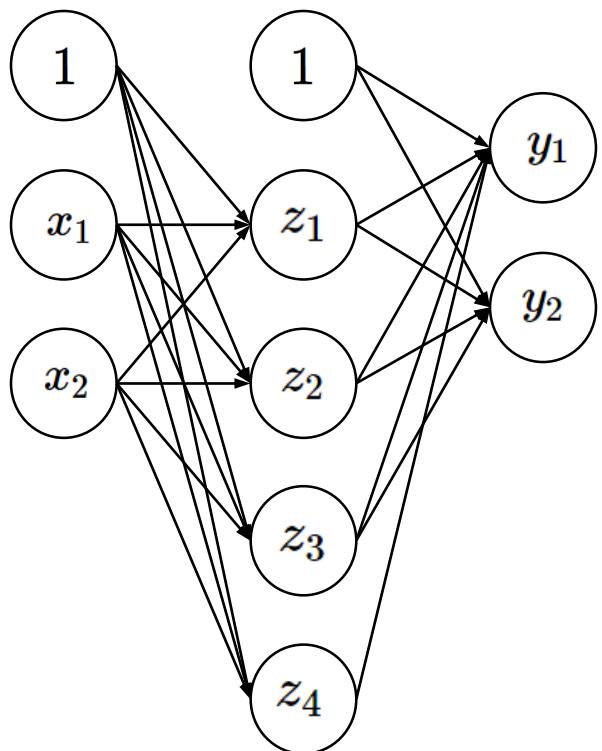
Do not include bias units



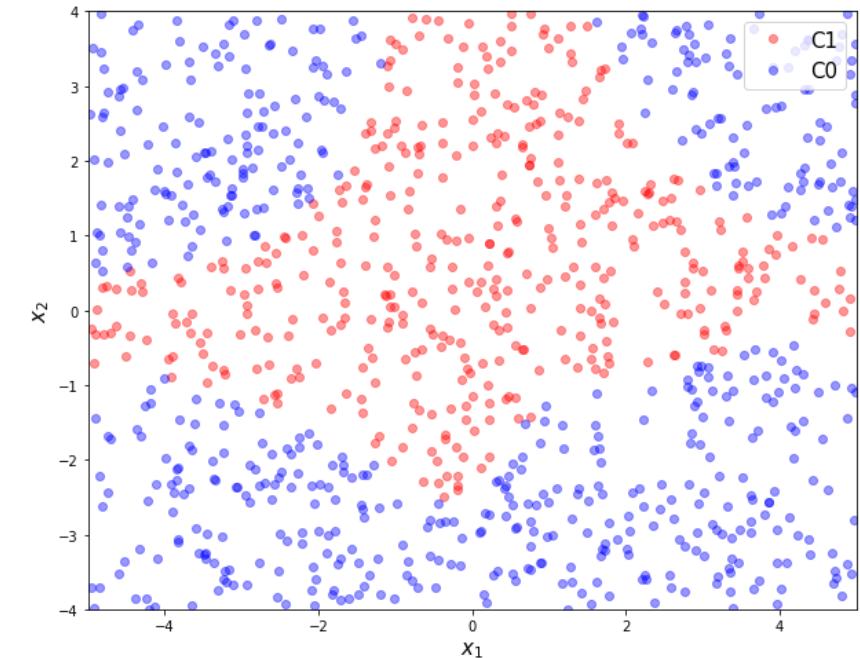
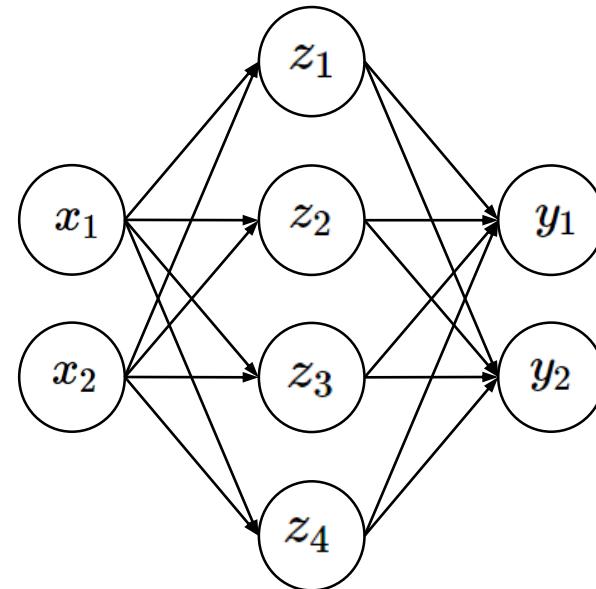
n_input = 2
n_hidden = 2
n_output = 2

Nonlinearly Distributed Data

- More neurons in hidden layer



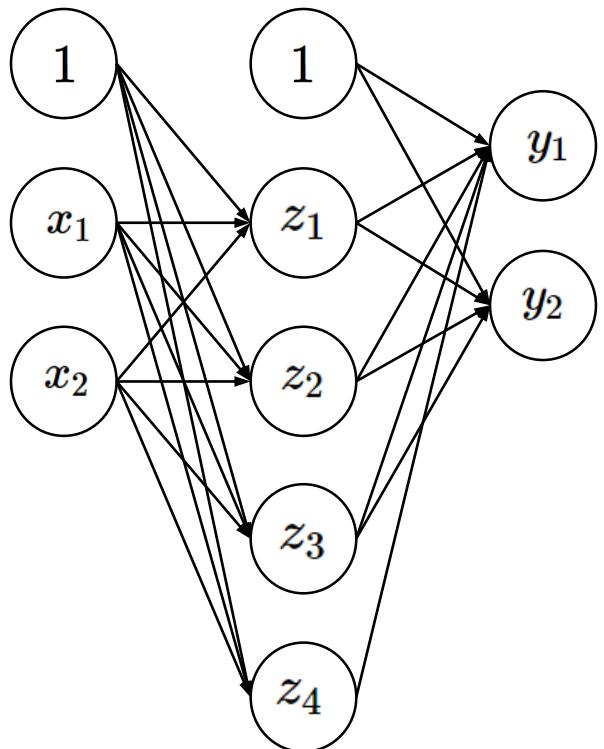
Do not include bias units



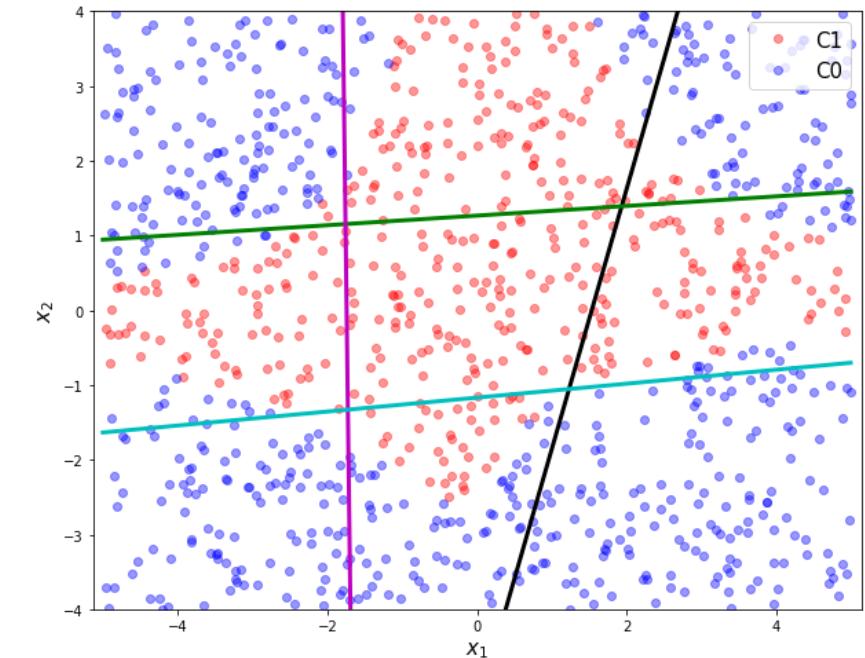
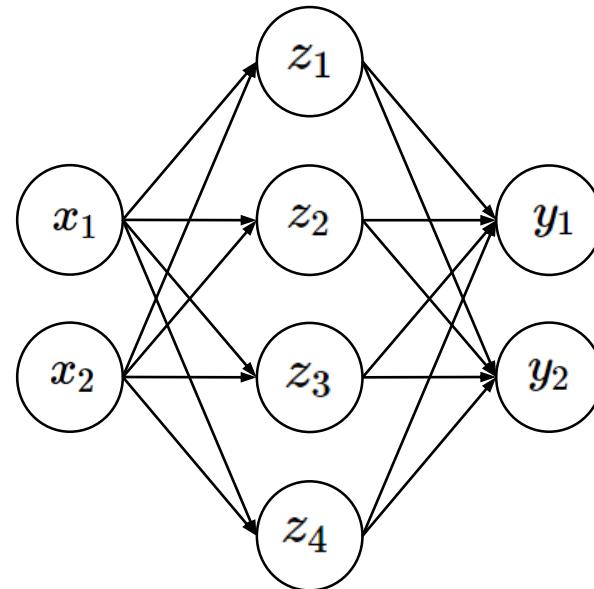
n_input = 2
n_hidden = 4
n_output = 2

Multi Layers

- Multiple linear classification boundaries



Do not include bias units



n_input = 2
n_hidden = 4
n_output = 2



(Artificial) Neural Networks: Training

Prof. Seungchul Lee
Industrial AI

Training Neural Networks: Optimization

- Learning or estimating weights and biases of multi-layer perceptron from training data
- 3 key components
 - objective function $f(\cdot)$
 - decision variable or unknown ω
 - constraints $g(\cdot)$
- In mathematical expression

$$\min_{\omega} f(\omega)$$

Training Neural Networks: Loss Function

- Measures error between target values and predictions

$$\min_{\omega} \sum_{i=1}^m \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$

- Example
 - Squared loss (for regression):
 - Cross entropy (for classification):

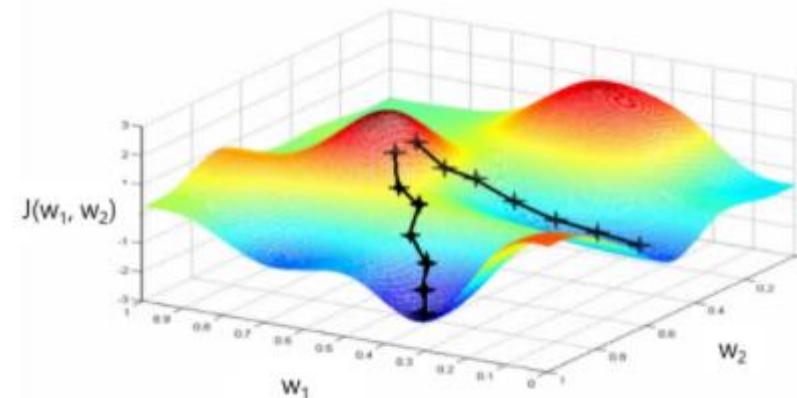
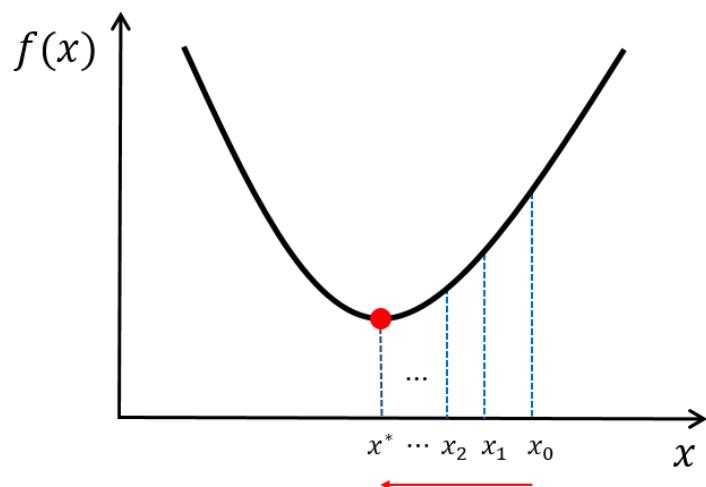
$$\frac{1}{m} \sum_{i=1}^m \left(h_{\omega} \left(x^{(i)} \right) - y^{(i)} \right)^2$$

$$-\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \left(h_{\omega} \left(x^{(i)} \right) \right) + \left(1 - y^{(i)} \right) \log \left(1 - h_{\omega} \left(x^{(i)} \right) \right)$$

Training Neural Networks: Gradient Descent

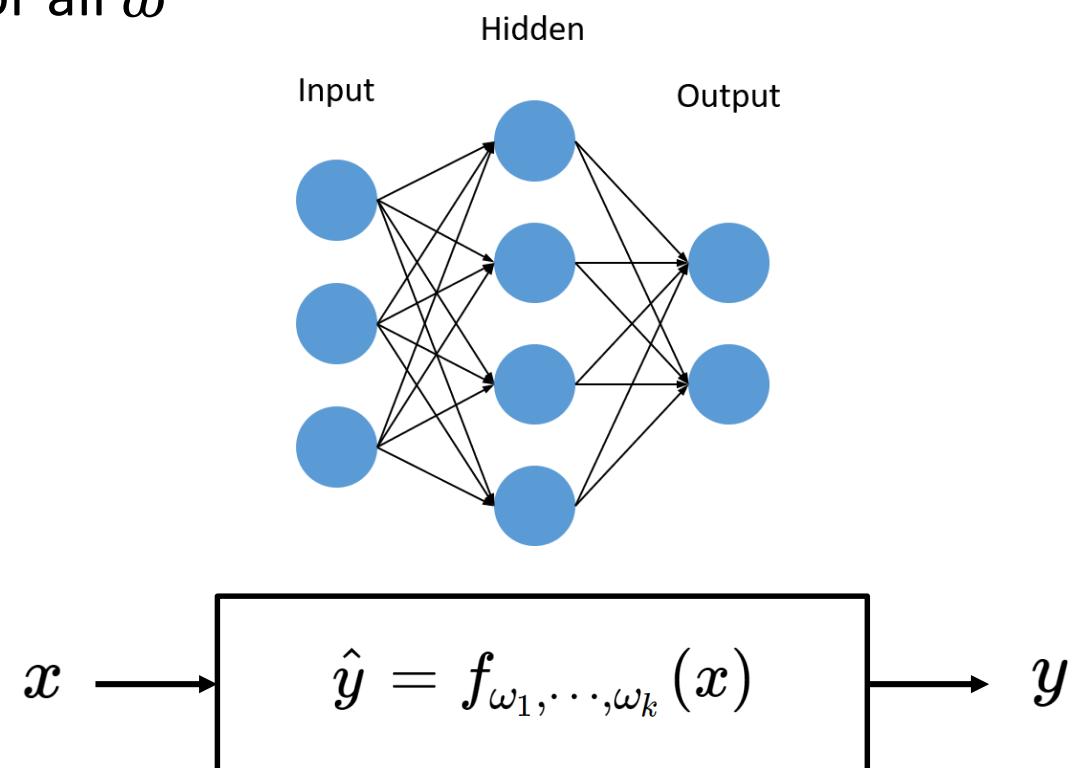
- Negative gradients points directly downhill of the cost function
- We can decrease the cost by moving in the direction of the negative gradient (α is a learning rate)

$$\omega \leftarrow \omega - \alpha \nabla_{\omega} \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$



Gradients in ANN

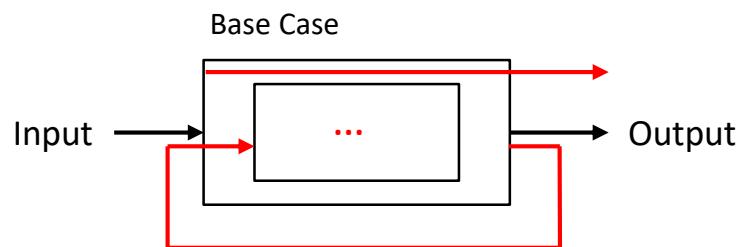
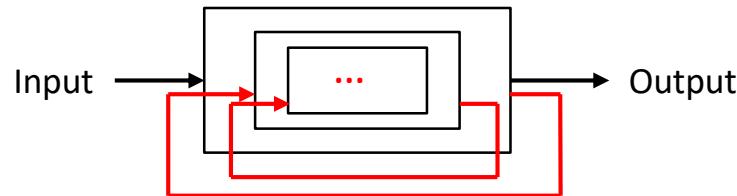
- Learning weights and biases from data using gradient descent
- $\frac{\partial \ell}{\partial \omega}$: too many computations are required for all ω
- Structural constraint of NN:
 - Composition of functions
 - Chain rule
 - Dynamic programming



Dynamic Programming

Recursive Algorithm

- One of the central ideas of computer science
- Depends on solutions to smaller instances of the same problem (= sub-problem)
- Function to call itself (it is impossible in the real world)
- Factorial example
 - $n! = n \cdot (n - 1) \cdots 2 \cdot 1$



Dynamic Programming

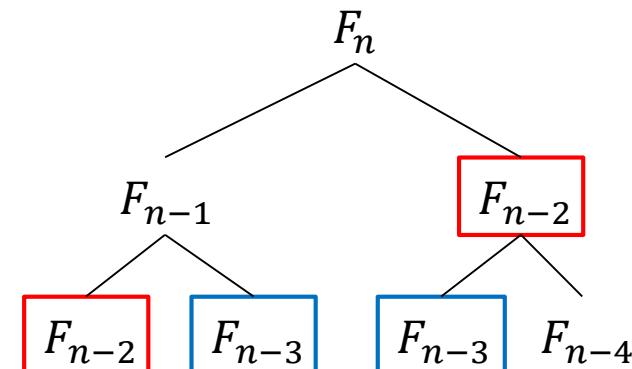
- Dynamic Programming: general, powerful algorithm design technique
- Fibonacci numbers:

$$\begin{aligned}F_1 &= F_2 = 1 \\F_n &= F_{n-1} + F_{n-2}\end{aligned}$$

Naïve Recursive Algorithm

```
fib(n) :  
    if  $n \leq 2$  :  $f = 1$   
    else :  $f = \text{fib}(n - 1) + \text{fib}(n - 2)$   
    return  $f$ 
```

- It works. Is it good?



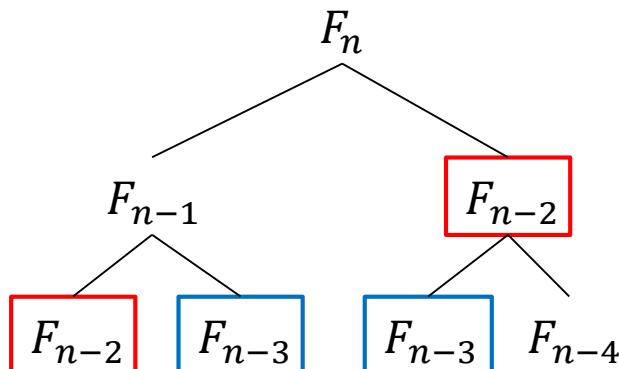
Memorized Recursive Algorithm

```
memo = [ ]
fib(n) :
    if n in memo : return memo[n]

    if n ≤ 2 : f = 1
    else : f = fib(n - 1) + fib(n - 2)

    memo[n] = f
    return f
```

- Benefit?
 - $\text{fib}(n)$ only recurses the first time it's called



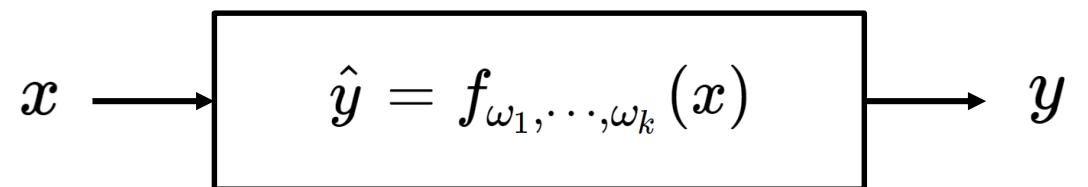
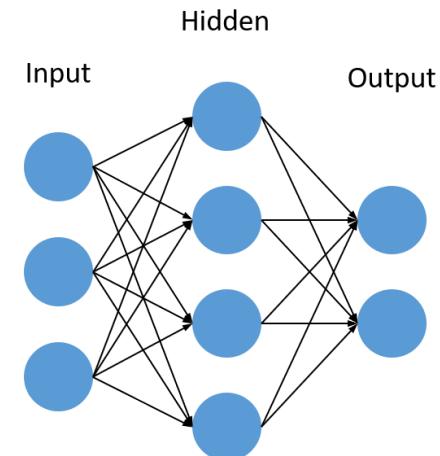
Dynamic Programming Algorithm

- Memorize (remember) & re-use solutions to subproblems that helps solve the problem
- DP \approx recursion + memorization

Backpropagation

Gradients in ANN

- Learning weights and biases from data using gradient descent
- $\frac{\partial \ell}{\partial \omega}$: too many computations are required for all ω
- Structural constraint of NN:
 - Composition of functions
 - Chain rule
 - Dynamic programming



Training Neural Networks: Backpropagation Learning

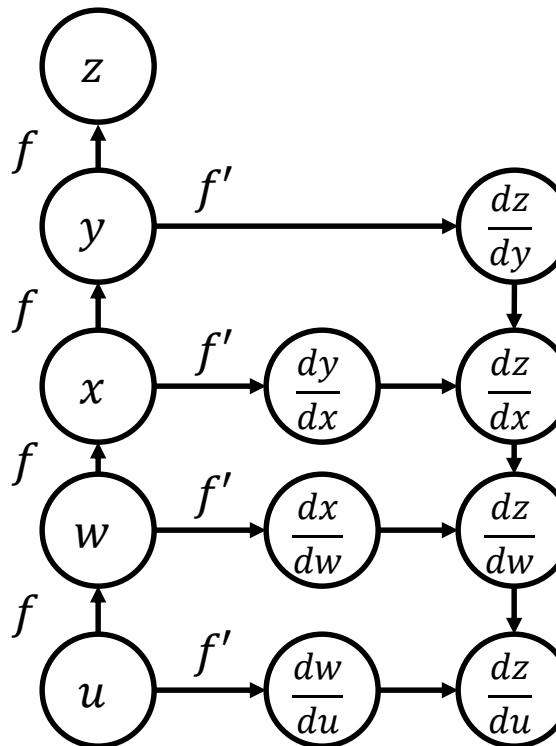
- Forward propagation
 - the initial information propagates up to the hidden units at each layer and finally produces output
- Backpropagation
 - allows the information from the cost to flow backwards through the network in order to compute the gradients

Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions

- $f(g(x))' = f'(g(x))g'(x)$
- $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
- $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \right) \cdot \frac{dx}{dw}$
- $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \right) \cdot \frac{dw}{du}$

- Backpropagation
 - Update weights recursively



Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions

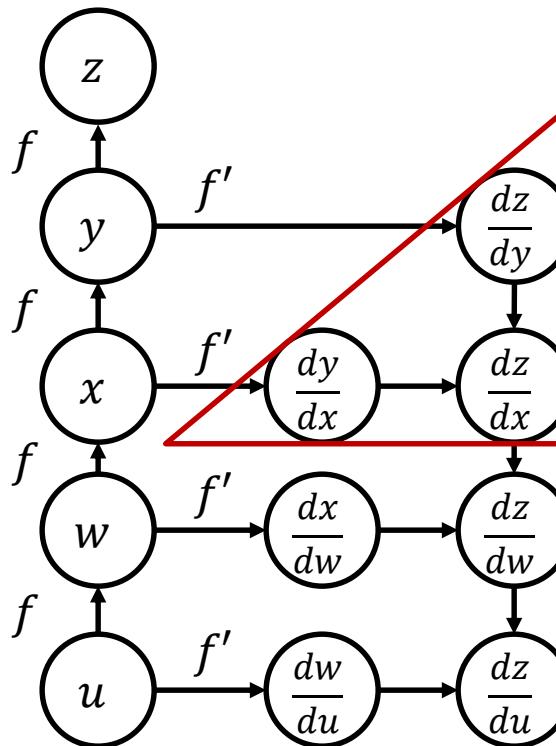
- $f(g(x))' = f'(g(x))g'(x)$

- $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$

- $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \right) \cdot \frac{dx}{dw}$

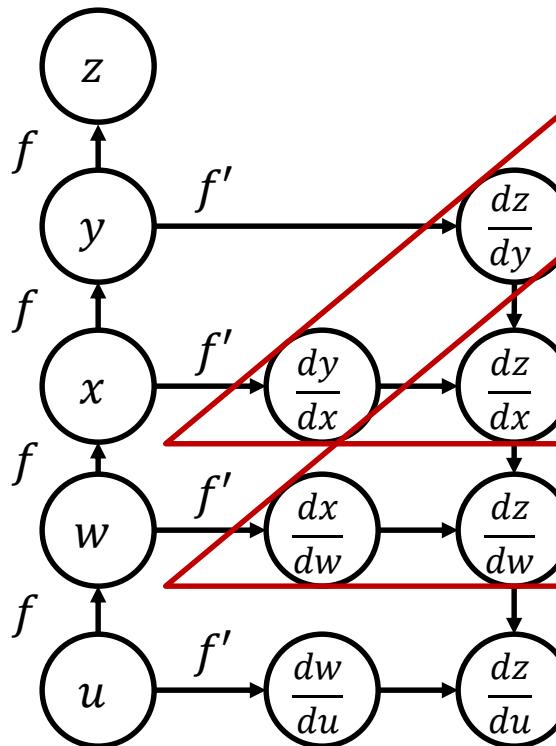
- $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \right) \cdot \frac{dw}{du}$

- Backpropagation
 - Update weights recursively



Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions
 - $f(g(x))' = f'(g(x))g'(x)$
 - $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
 - $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \right) \cdot \frac{dx}{dw}$
 - $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \right) \cdot \frac{dw}{du}$
 - Backpropagation
 - Update weights recursively



Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions

- $f(g(x))' = f'(g(x))g'(x)$

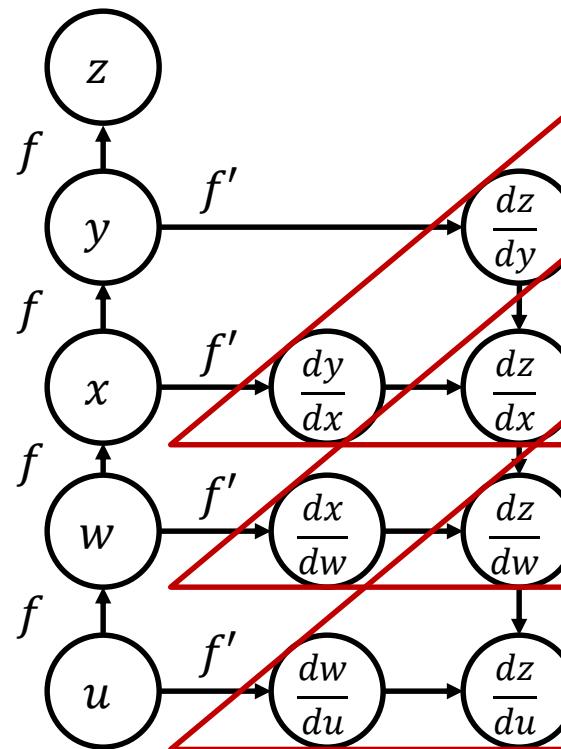
- $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$

- $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \right) \cdot \frac{dx}{dw}$

- $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \right) \cdot \frac{dw}{du}$

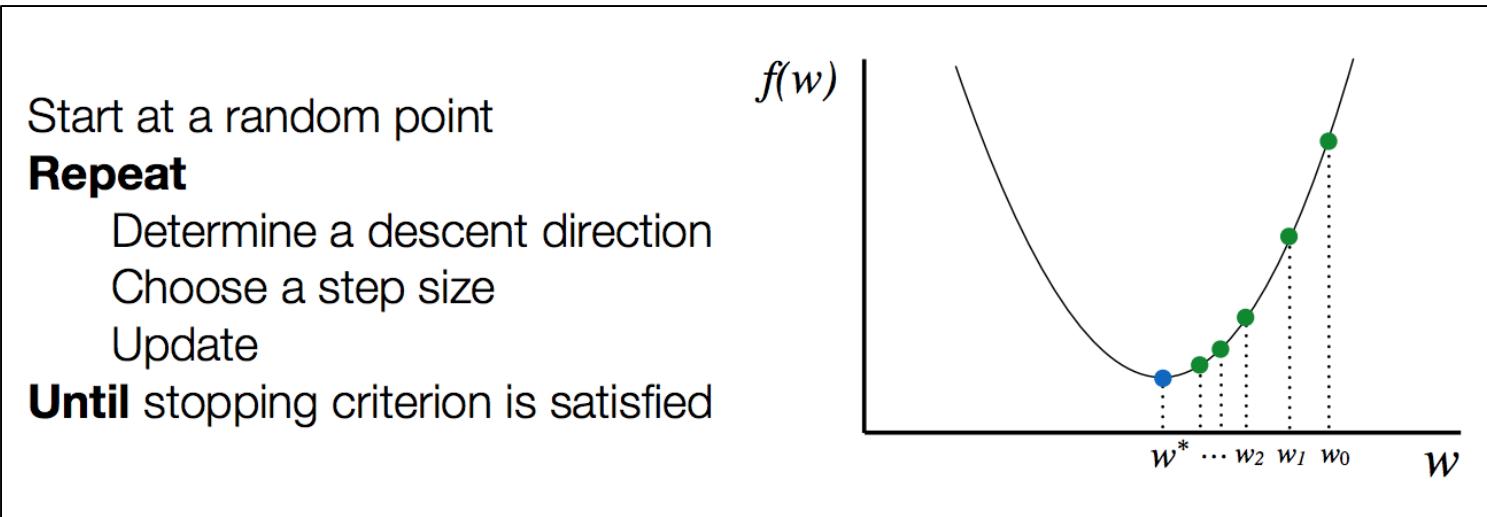
- Backpropagation

- Update weights recursively with memory



Training Neural Networks with TensorFlow

- Optimization procedure

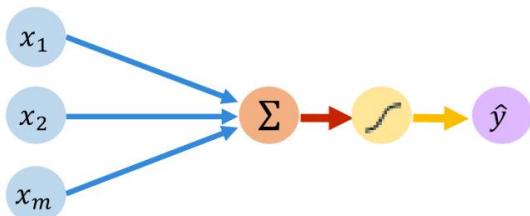


- It is not easy to numerically compute gradients in network in general.
 - The good news: people have already done all the "hard work" of developing numerical solvers (or libraries)
 - There are a wide range of tools → We will use the TensorFlow

Core Foundation Review

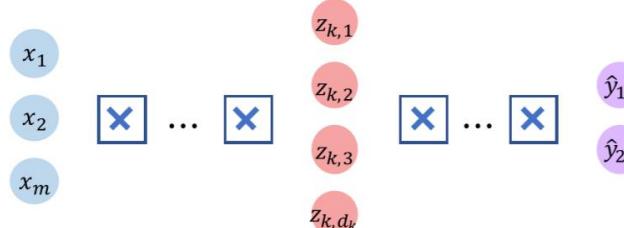
The Perceptron

- Structural building blocks
- Nonlinear activation functions



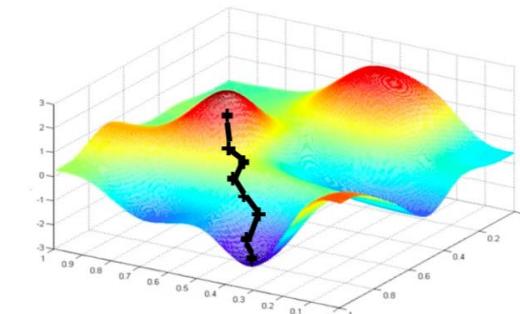
Neural Networks

- Stacking Perceptrons to form neural networks
- Optimization through backpropagation



Training in Practice

- Adaptive learning
- Batching
- Regularization



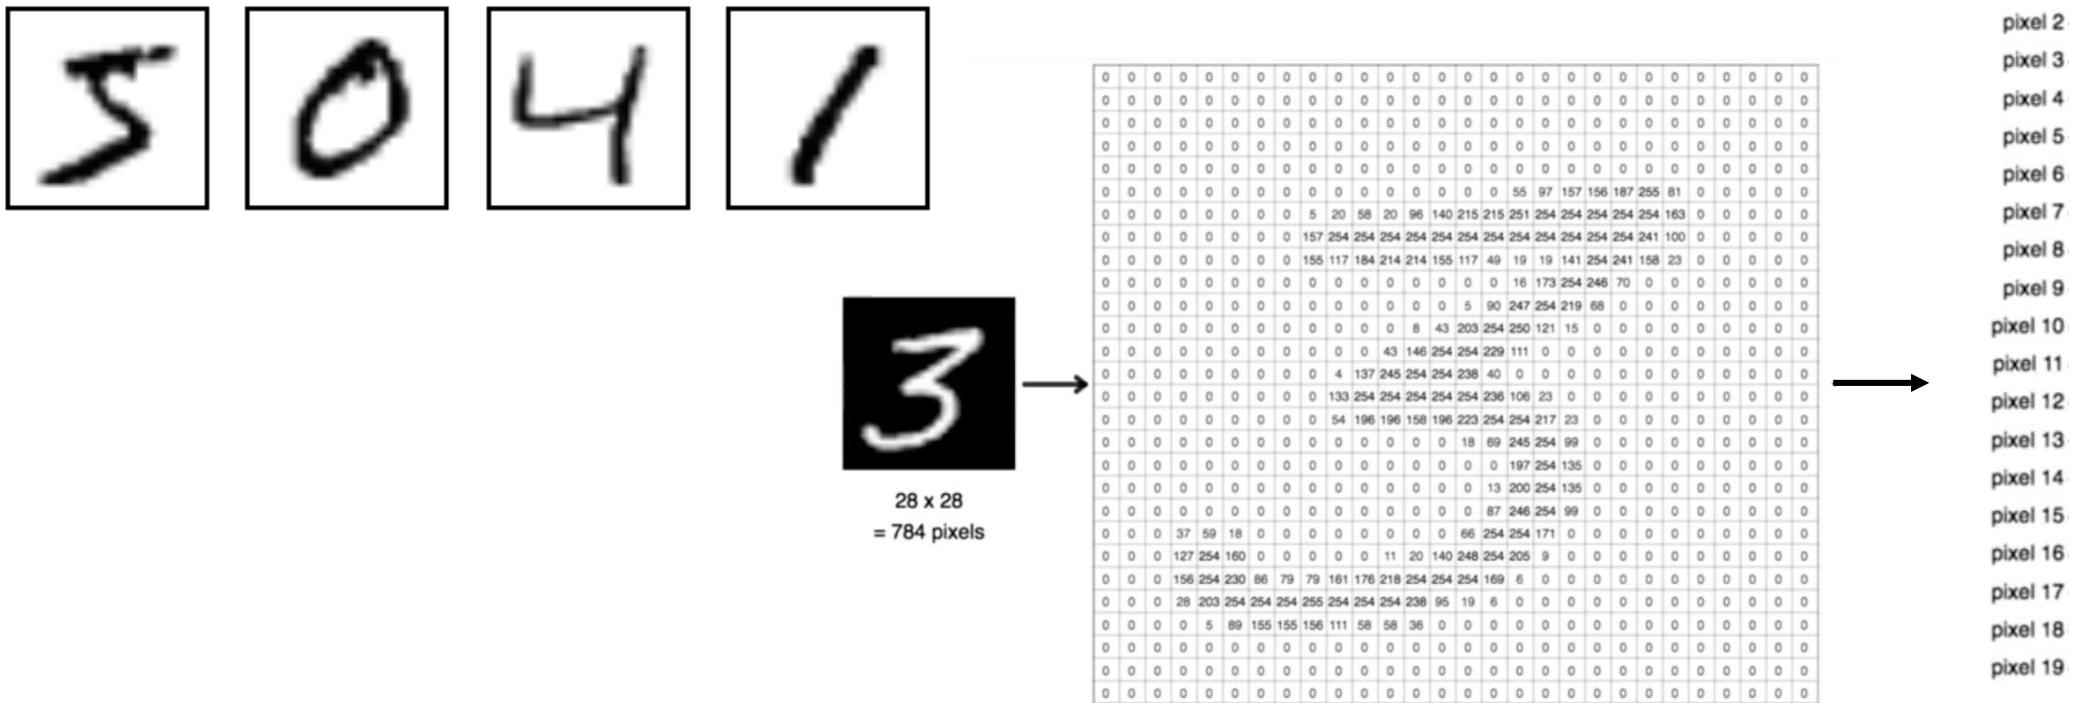


(Artificial) Neural Networks with TensorFlow

Prof. Seungchul Lee
Industrial AI Lab.

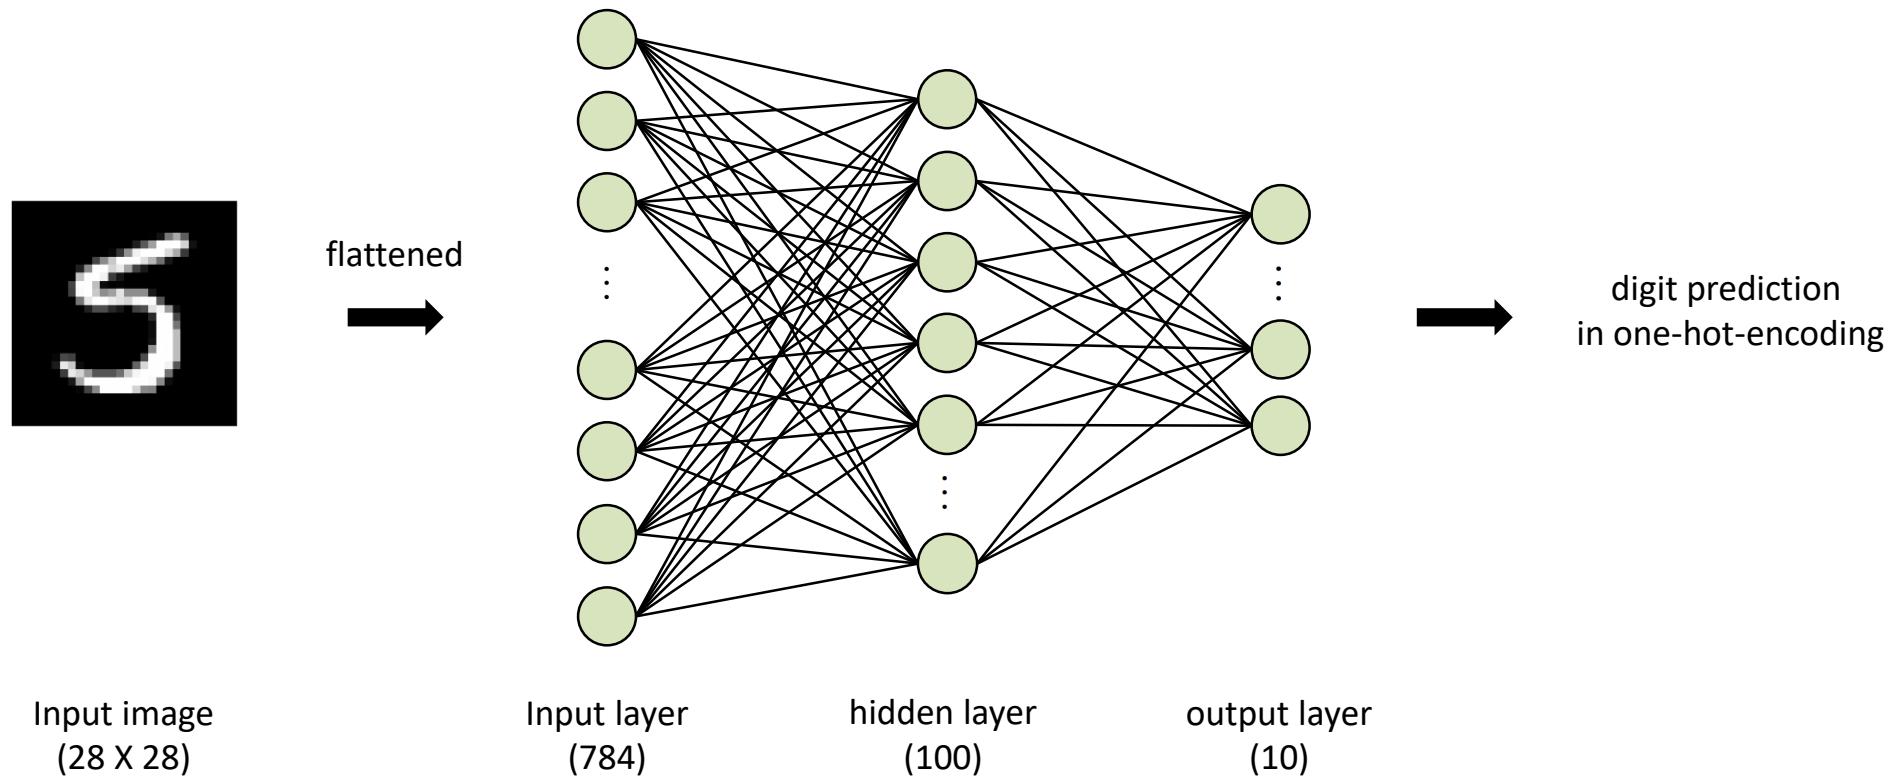
MNIST database

- Mixed National Institute of Standards and Technology database
 - Handwritten digit database
 - 28×28 gray scaled image
 - **Flattened** matrix into a vector of $28 \times 28 = 784$



ANN in TensorFlow: MNIST

Our Network Model

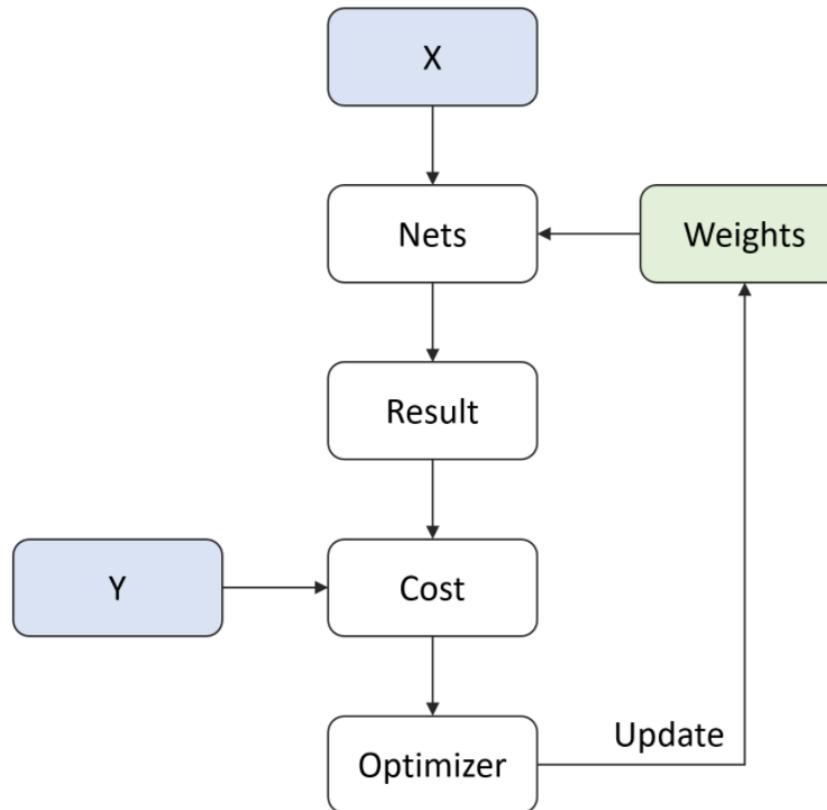


Iterative Optimization

- We will use
 - Mini-batch gradient descent
 - Adam optimizer

$$\begin{aligned} \min_{\theta} \quad & f(\theta) \\ \text{subject to} \quad & g_i(\theta) \leq 0 \end{aligned}$$

$$\theta := \theta - \alpha \nabla_{\theta} \left(h_{\theta} \left(x^{(i)} \right), y^{(i)} \right)$$



ANN with TensorFlow

- Import Library

```
# Import Library
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
```

- Load MNIST Data
 - Download MNIST data from TensorFlow tutorial example

```
from tensorflow.examples.tutorials.mnist import input_data

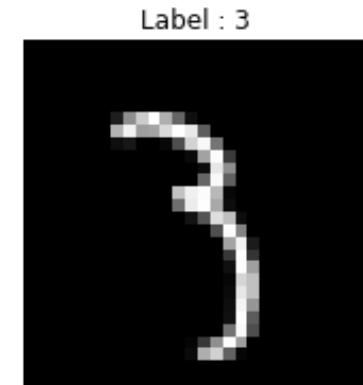
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

One Hot Encoding

- Batch maker

```
train_x, train_y = mnist.train.next_batch(1)
img = train_x[0,:].reshape(28,28)

plt.figure(figsize=(5,3))
plt.imshow(img,'gray')
plt.title("Label : {}".format(np.argmax(train_y[0,:])))
plt.xticks([])
plt.yticks([])
plt.show()
```



- One hot encoding

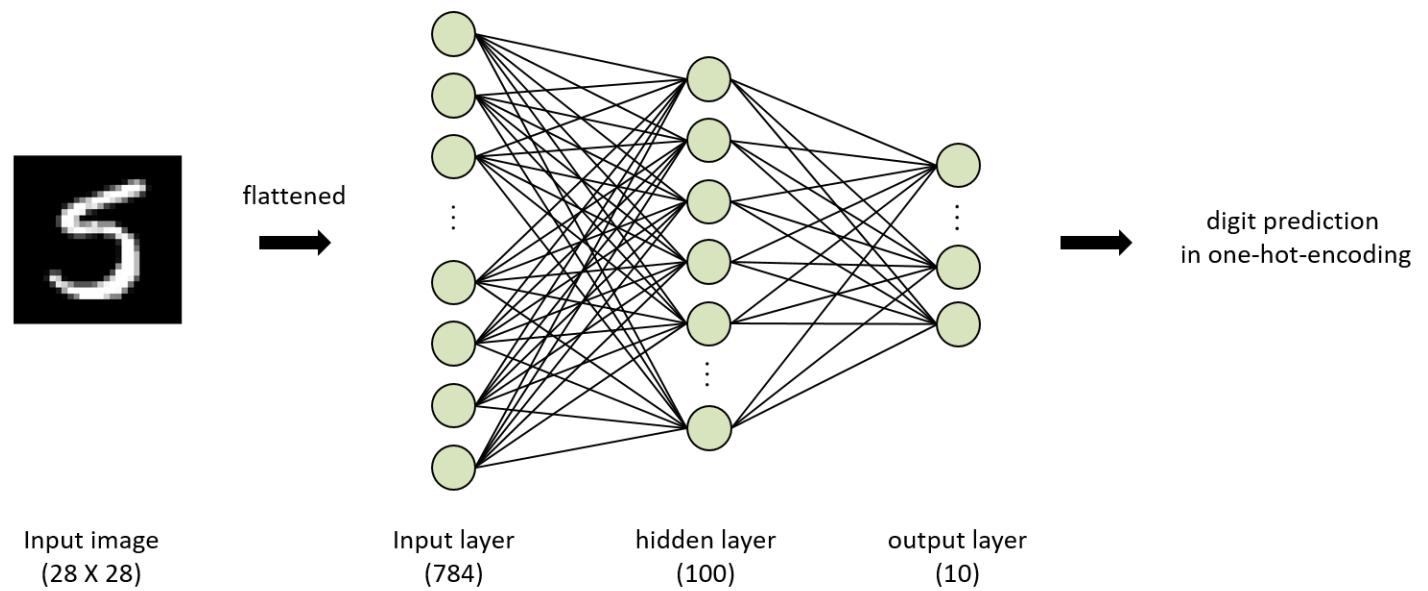
```
print ('Train labels : {}'.format(train_y[0, :]))
```

Train labels : [0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]

ANN Structure

- Input size
- Hidden layer size
- The number of classes

```
n_input = 28*28  
n_hidden = 100  
n_output = 10
```



Weights & Biases and Placeholder

- Define parameters based on predefined layer size
- Initialize with normal distribution with $\mu = 0$ and $\sigma = 0.1$

```
weights = {
    'hidden' : tf.Variable(tf.random_normal([n_input, n_hidden], stddev = 0.1)),
    'output' : tf.Variable(tf.random_normal([n_hidden, n_output], stddev = 0.1))
}

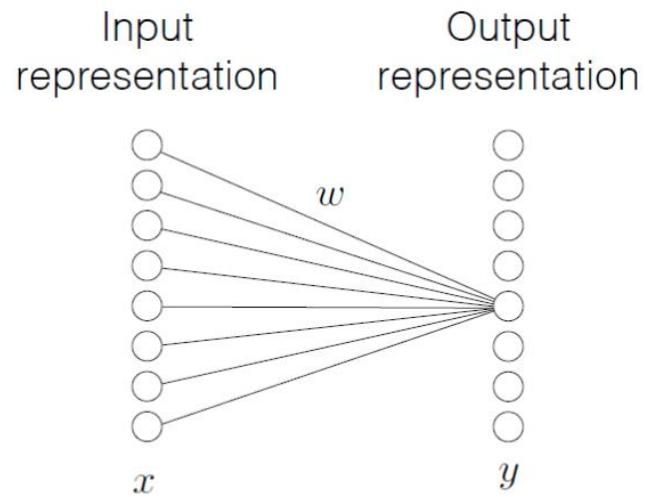
biases = {
    'hidden' : tf.Variable(tf.random_normal([n_hidden], stddev = 0.1)),
    'output' : tf.Variable(tf.random_normal([n_output], stddev = 0.1))
}
```

- Placeholder

```
x = tf.placeholder(tf.float32, [None, n_input])
y = tf.placeholder(tf.float32, [None, n_output])
```

Build a Model

- First, the layer performs several matrix multiplication to produce a set of linear activations



$$y_j = \left(\sum_i \omega_{ij} x_i \right) + b_j$$
$$y = \omega^T x + b$$

```
# Define Network
def build_model(x, weights, biases):

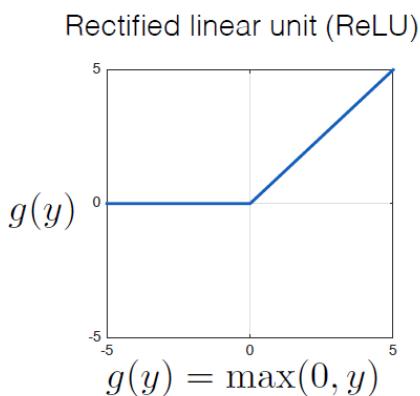
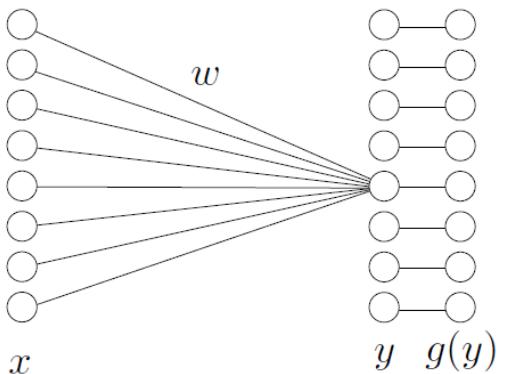
    # first hidden layer
    hidden = tf.add(tf.matmul(x, weights['hidden']), biases['hidden'])
    # non-linear activate function
    hidden = tf.nn.relu(hidden)

    # Output layer
    output = tf.add(tf.matmul(hidden, weights['output']), biases['output'])

    return output
```

Build a Model

- Second, each linear activation is running through a nonlinear activation function



```
# Define Network
def build_model(x, weights, biases):

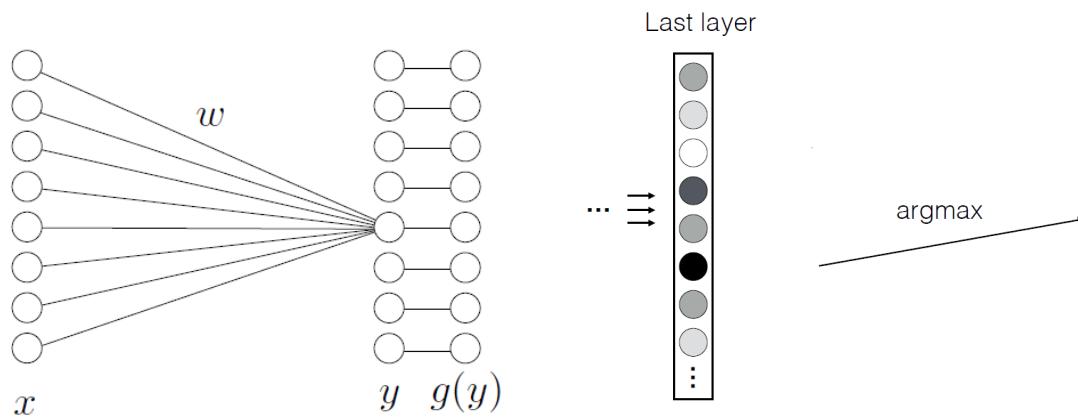
    # first hidden layer
    hidden = tf.add(tf.matmul(x, weights['hidden']), biases['hidden'])
    # non-linear activate function
    hidden = tf.nn.relu(hidden)

    # Output layer
    output = tf.add(tf.matmul(hidden, weights['output']), biases['output'])

    return output
```

Build a Model

- Third, predict values with an affine transformation



```
# Define Network
def build_model(x, weights, biases):

    # first hidden layer
    hidden = tf.add(tf.matmul(x, weights['hidden']), biases['hidden'])
    # non-linear activate function
    hidden = tf.nn.relu(hidden)

    # Output layer
    output = tf.add(tf.matmul(hidden, weights['output']), biases['output'])

    return output
```

Loss and Optimizer

- Loss: softmax cross entropy

$$-\frac{1}{N} \sum_{i=1}^N y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))$$

- Optimizer

- AdamOptimizer: the most popular optimizer

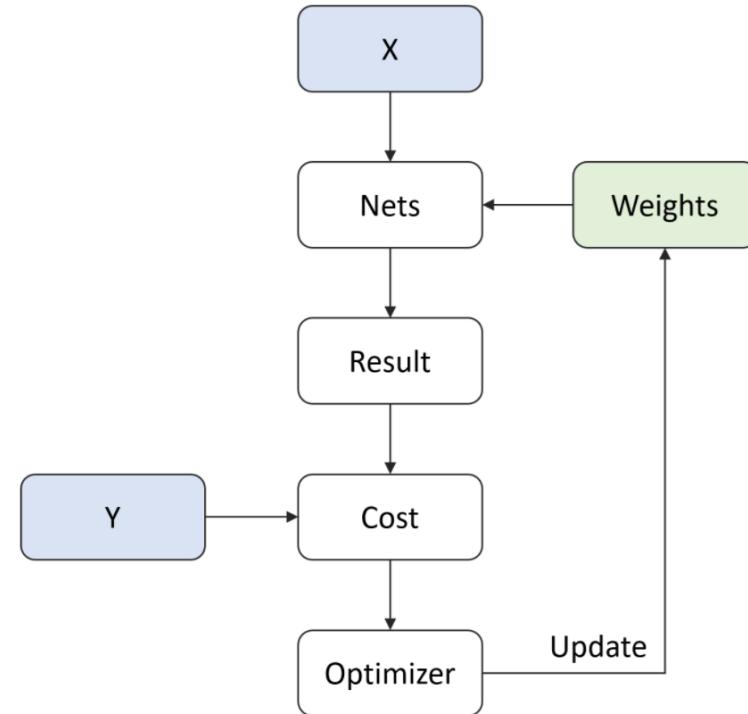
```
# Define Loss
pred = build_model(x, weights, biases)
loss = tf.nn.softmax_cross_entropy_with_logits(logits = pred, labels = y)
loss = tf.reduce_mean(loss)

LR = 0.0001
optm = tf.train.AdamOptimizer(LR).minimize(loss)
```

Iteration Configuration

- Define parameters for training ANN
 - n_batch: batch size for mini-batch gradient descent
 - n_iter: the number of iteration steps
 - n_prt: check loss for every n_prt iteration

```
n_batch = 50      # Batch Size
n_iter = 5000    # Learning Iteration
n_prt = 250      # Print Cycle
```



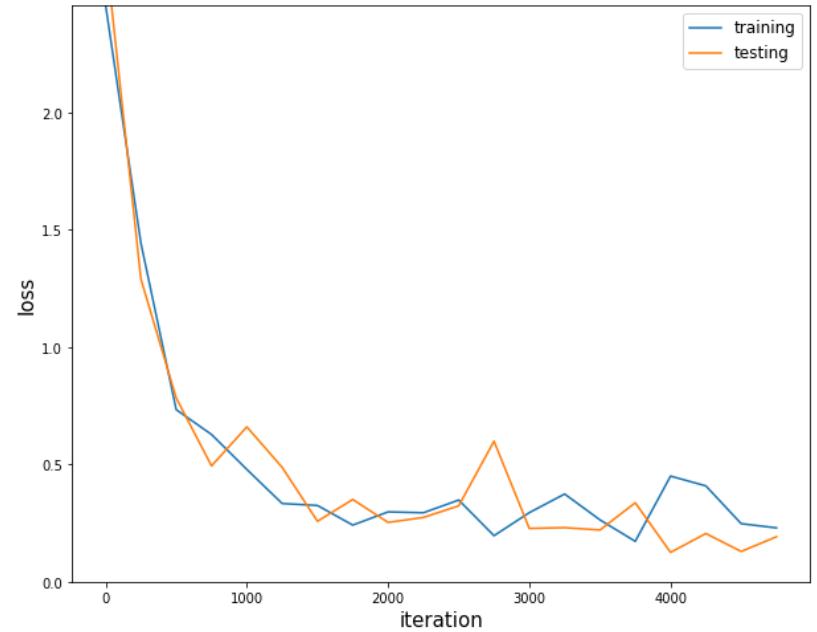
Optimization

```
sess = tf.Session()
init = tf.global_variables_initializer()
sess.run(init)

loss_record_training = []
loss_record_testing = []
for epoch in range(n_iter):
    train_x, train_y = mnist.train.next_batch(n_batch)
    sess.run(optm, feed_dict = {x: train_x, y: train_y})

    if epoch % n_prt == 0:
        test_x, test_y = mnist.test.next_batch(n_batch)
        c1 = sess.run(loss, feed_dict = {x: train_x, y: train_y})
        c2 = sess.run(loss, feed_dict = {x: test_x, y: test_y})
        loss_record_training.append(c1)
        loss_record_testing.append(c2)
        print ("Iter : {}".format(epoch))
        print ("Cost : {}".format(c1))

plt.figure(figsize=(10,8))
plt.plot(np.arange(len(loss_record_training))*n_prt,
         loss_record_training, label = 'training')
plt.plot(np.arange(len(loss_record_testing))*n_prt,
         loss_record_testing, label = 'testing')
plt.xlabel('iteration', fontsize = 15)
plt.ylabel('loss', fontsize = 15)
plt.legend(fontsize = 12)
plt.ylim([0,np.max(loss_record_training)])
plt.show()
```



Test or Evaluation

```
test_x, test_y = mnist.test.next_batch(100)
my_pred = sess.run(pred, feed_dict = {x : test_x})
my_pred = np.argmax(my_pred, axis = 1)

labels = np.argmax(test_y, axis = 1)

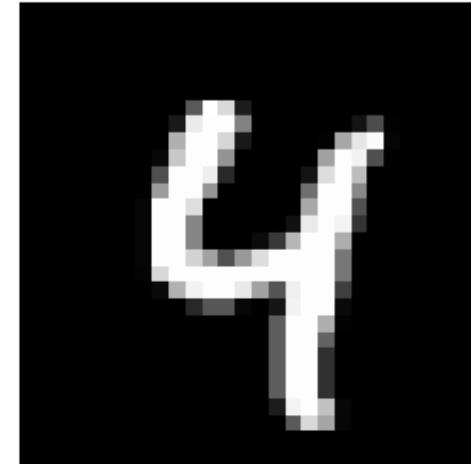
accr = np.mean(np.equal(my_pred, labels))
print("Accuracy : {}%".format(accr*100))
```

Accuracy : 96.0%

```
test_x, test_y = mnist.test.next_batch(1)
logits = sess.run(tf.nn.softmax(pred), feed_dict = {x : test_x})
predict = np.argmax(logits)

plt.imshow(test_x.reshape(28,28), 'gray')
plt.xticks([])
plt.yticks([])
plt.show()

print('Prediction : {}'.format(predict))
np.set_printoptions(precision=2, suppress=True)
print('Probability : {}'.format(logits.ravel()))
```



Prediction : 4
Probability : [0. 0. 0. 0. 0.9 0. 0. 0.01 0. 0.09]



(Artificial) Neural Networks: Advanced

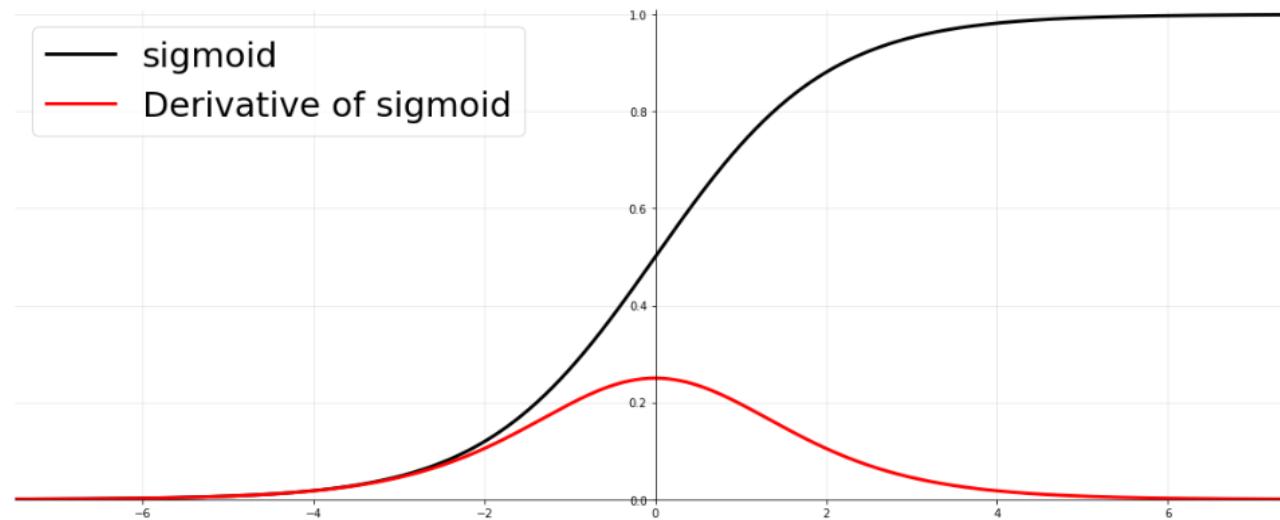
**Prof. Seungchul Lee
Industrial AI Lab.**

Nonlinear Activation Function

The Vanishing Gradient Problem

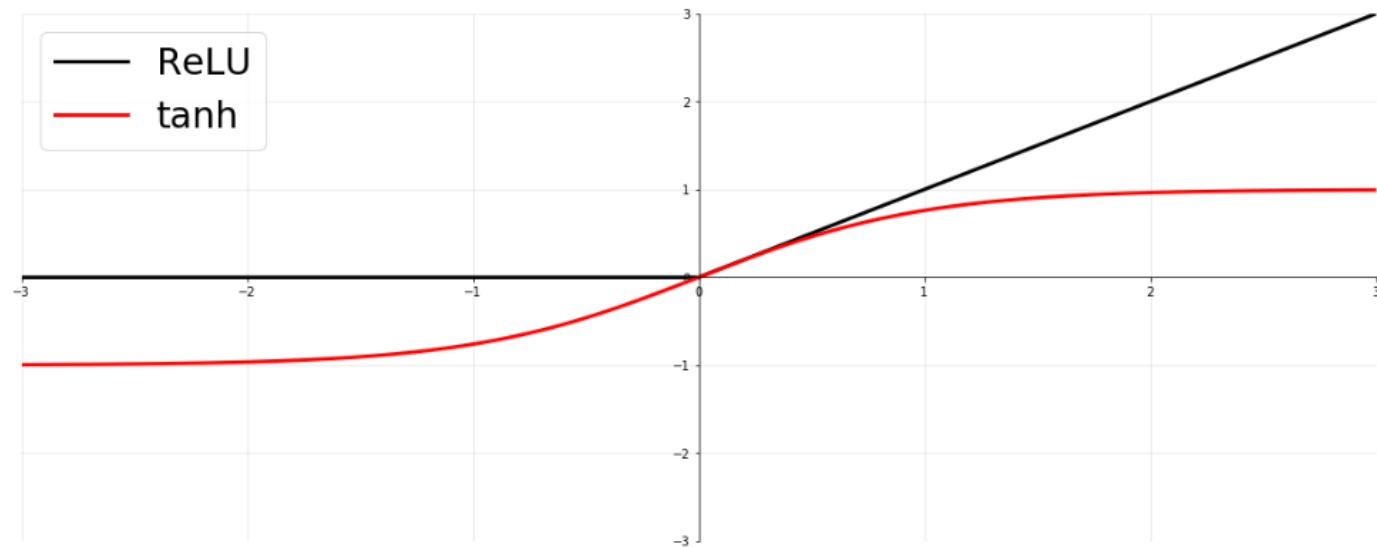
- As more layers using certain activation functions are added to neural networks, the gradients of the loss function approaches zero, making the network hard to train.
- For example,

$$-\frac{dz}{du} = \frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \cdot \frac{dw}{du}$$



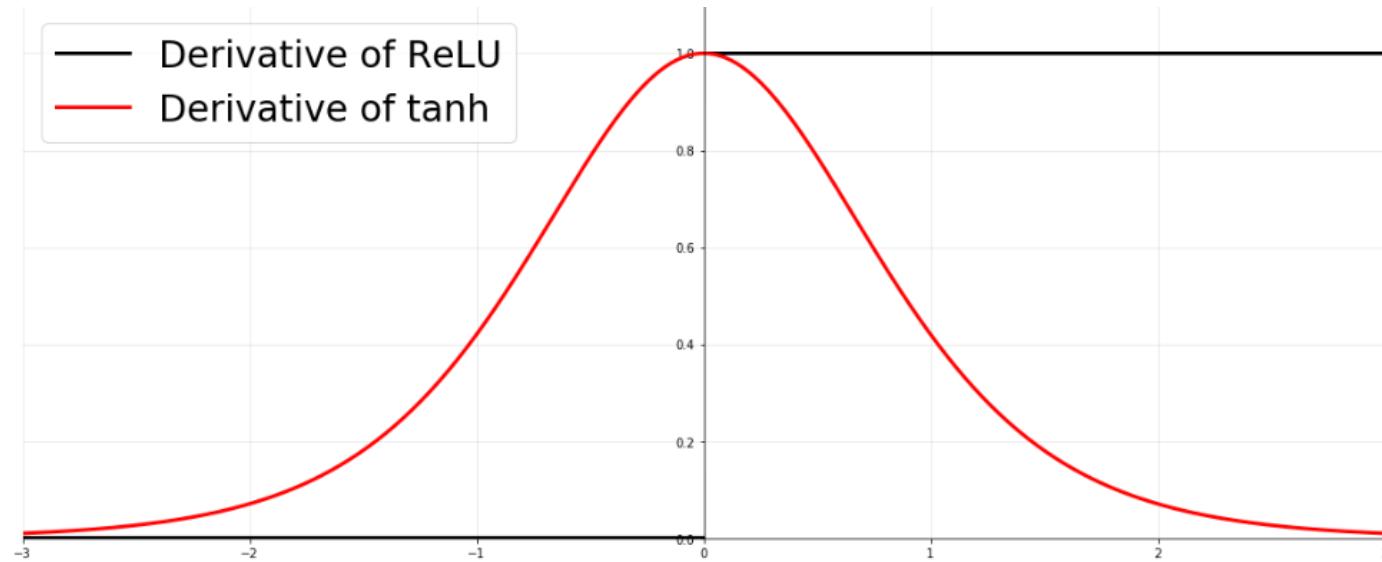
Rectifiers

- The use of the **ReLU** activation function was a great improvement compared to the historical tanh.



Rectifiers

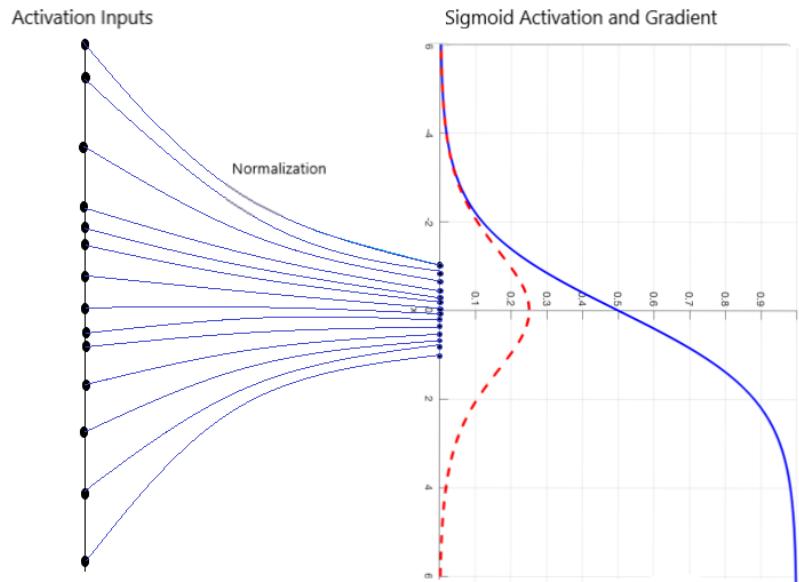
- This can be explained by the derivative of ReLU itself not vanishing, and by the resulting coding being sparse (Glorot et al., 2011).



Batch Normalization

Batch Normalization

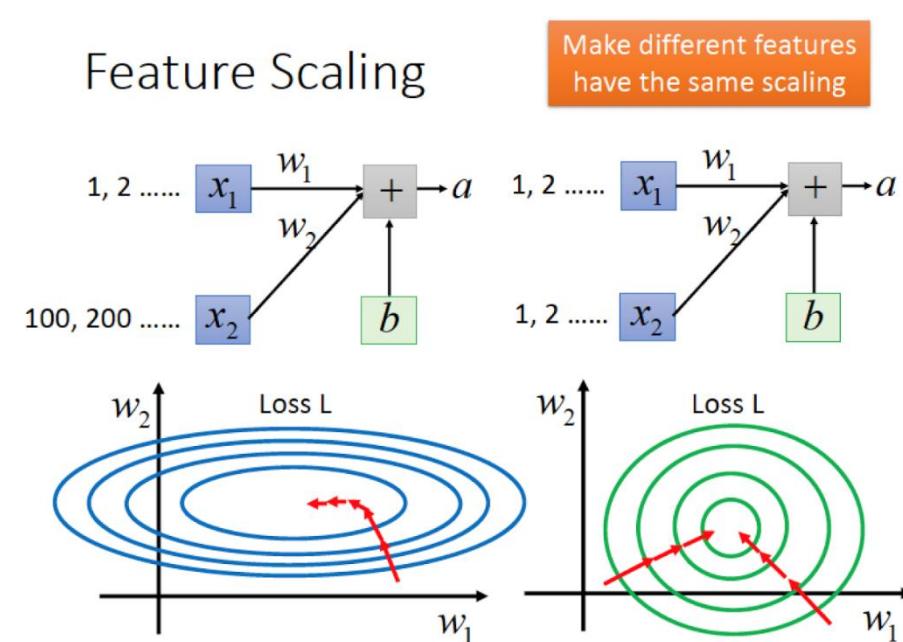
- Batch normalization is a technique for improving the performance and stability of artificial neural networks.
- It is used to normalize the input layer by adjusting and scaling the activations.



S. Ioffe and C. Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In International Conference on Machine Learning (ICML), 2015.

Batch Normalization

- Batch normalization is a technique for improving the performance and stability of artificial neural networks.
- It is used to normalize the input layer by adjusting and scaling the activations.



S. Ioffe and C. Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In International Conference on Machine Learning (ICML), 2015.

Batch Normalization

- During training batch normalization shifts and rescales according to the mean and variance estimated on the batch.
- During test, it simply shifts and rescales according to the empirical moments estimated during training.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Dropout as Regularization

Regularization (Shrinkage Methods)

- Often, overfitting associated with very large estimated parameters
- We want to balance
 - how well function fits data
 - magnitude of coefficients

$$\text{Total cost} = \underbrace{\text{measure of fit}}_{RSS(\theta)} + \lambda \cdot \underbrace{\text{measure of magnitude of coefficients}}_{\lambda \cdot \|\theta\|_2^2}$$

$$\implies \min \|\Phi\theta - y\|_2^2 + \lambda \|\theta\|_2^2$$

- multi-objective optimization
- λ is a tuning parameter

Different Regularization Techniques

- Big Data
- Data augmentation
 - The simplest way to reduce overfitting is to increase the size of the training data.

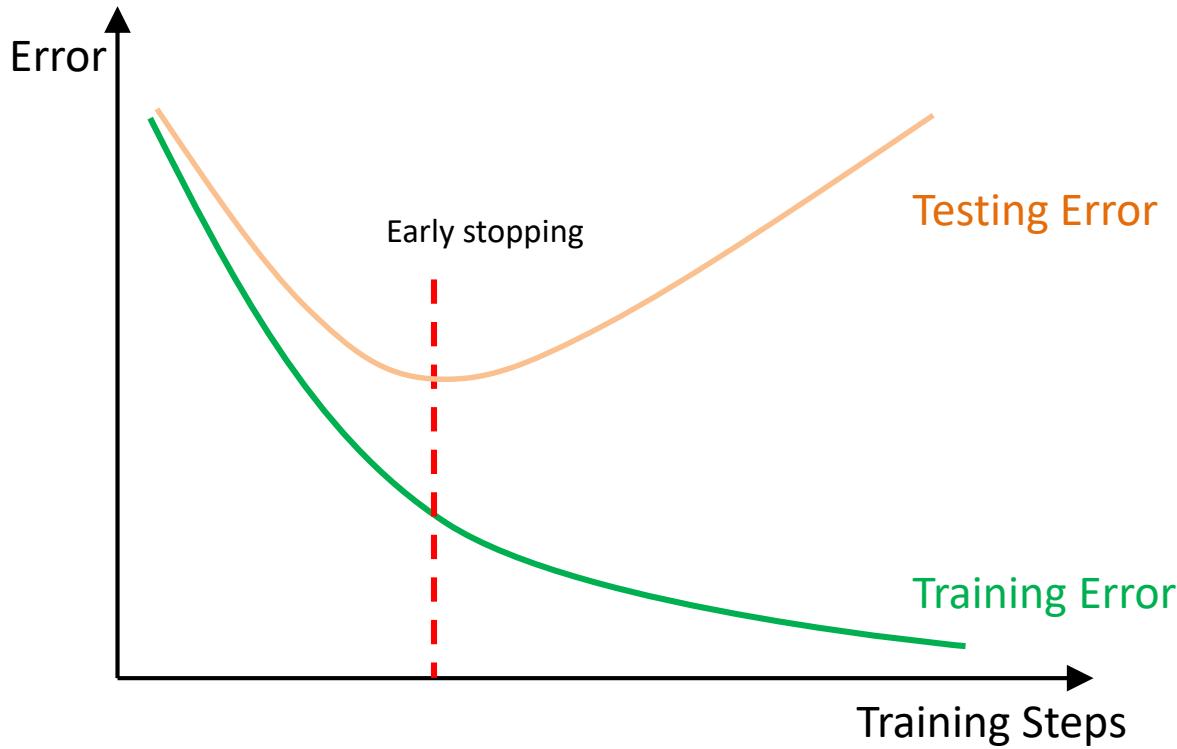


shift shift shear shift & scale rotate & scale



Different Regularization Techniques

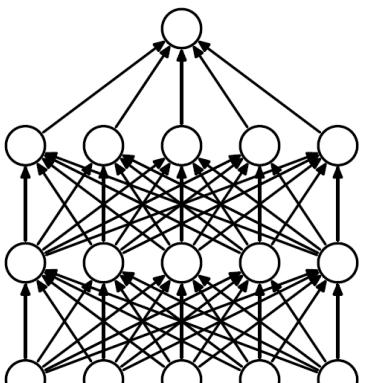
- Early stopping
 - When we see that the performance on the validation set is getting worse, we immediately stop the training on the model.



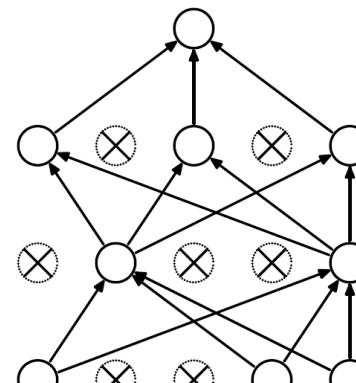
Different Regularization Techniques in Deep Learning

- **Dropout**

- This is the one of the most interesting types of regularization techniques.
- It also produces very good results and is consequently the most frequently used regularization technique in the field of deep learning.
- At every iteration, it **randomly selects some nodes and removes them**.
- It can also be thought of **as an ensemble technique** in machine learning.



(a) Standard Neural Net

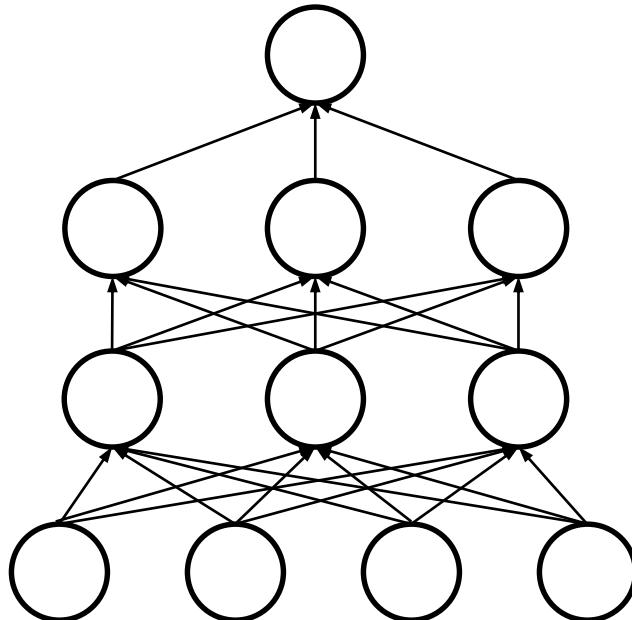


(b) After applying dropout.

Dropout Illustration

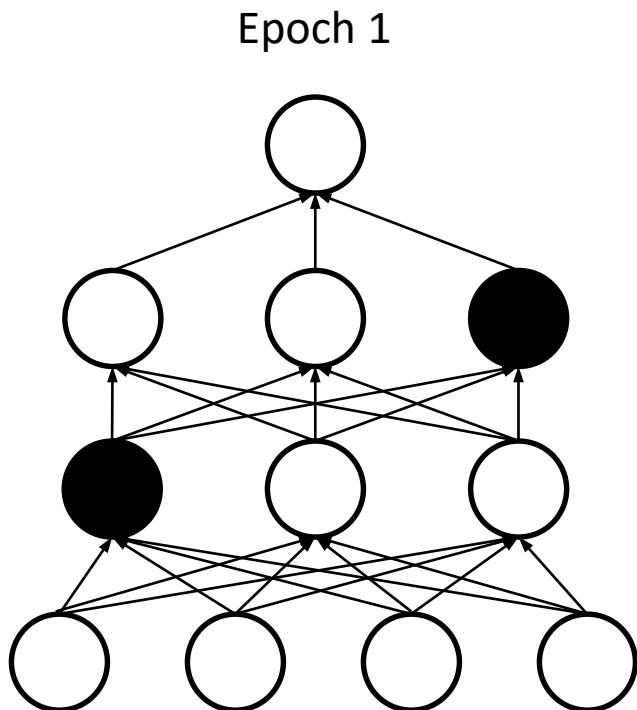
- Effectively, a different architecture at every training epoch
- It can also be thought of as an ensemble technique in machine learning.

Original model



Dropout Illustration

- Effectively, a different architecture at every training epoch
- It can also be thought of as an ensemble technique in machine learning.

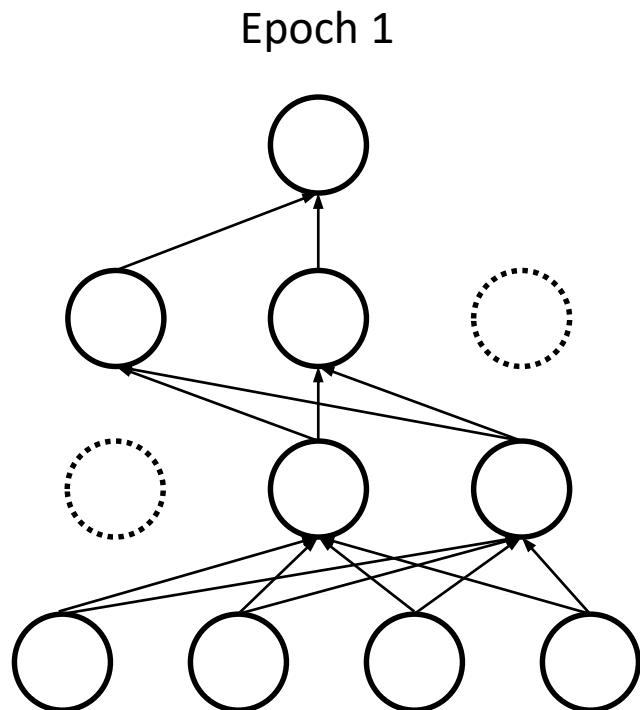


`tf.nn.dropout(layer, rate = p)`

rate: the probability that each element is dropped. For example, setting rate = 0.1 would drop 10% of input elements

Dropout Illustration

- Effectively, a different architecture at every training epoch
- It can also be thought of as an ensemble technique in machine learning.



`tf.nn.dropout(layer, rate = p)`

rate: the probability that each element is dropped. For example, setting rate = 0.1 would drop 10% of input elements

Dropout

- For training
 - rate: the probability that each element is dropped. For example, setting rate = 0.1 would drop 10% of input elements
 - With probability rate, drops elements of layers. Input that are kept are scaled up by $\frac{1}{1-\text{rate}}$, otherwise outputs 0. The scaling is so that the expected sum is unchanged.

- For testing
 - All the elements are kept

